

Lifelong Learning

CS 330

Plan for Today

The lifelong learning **problem statement**

Basic approaches to lifelong learning

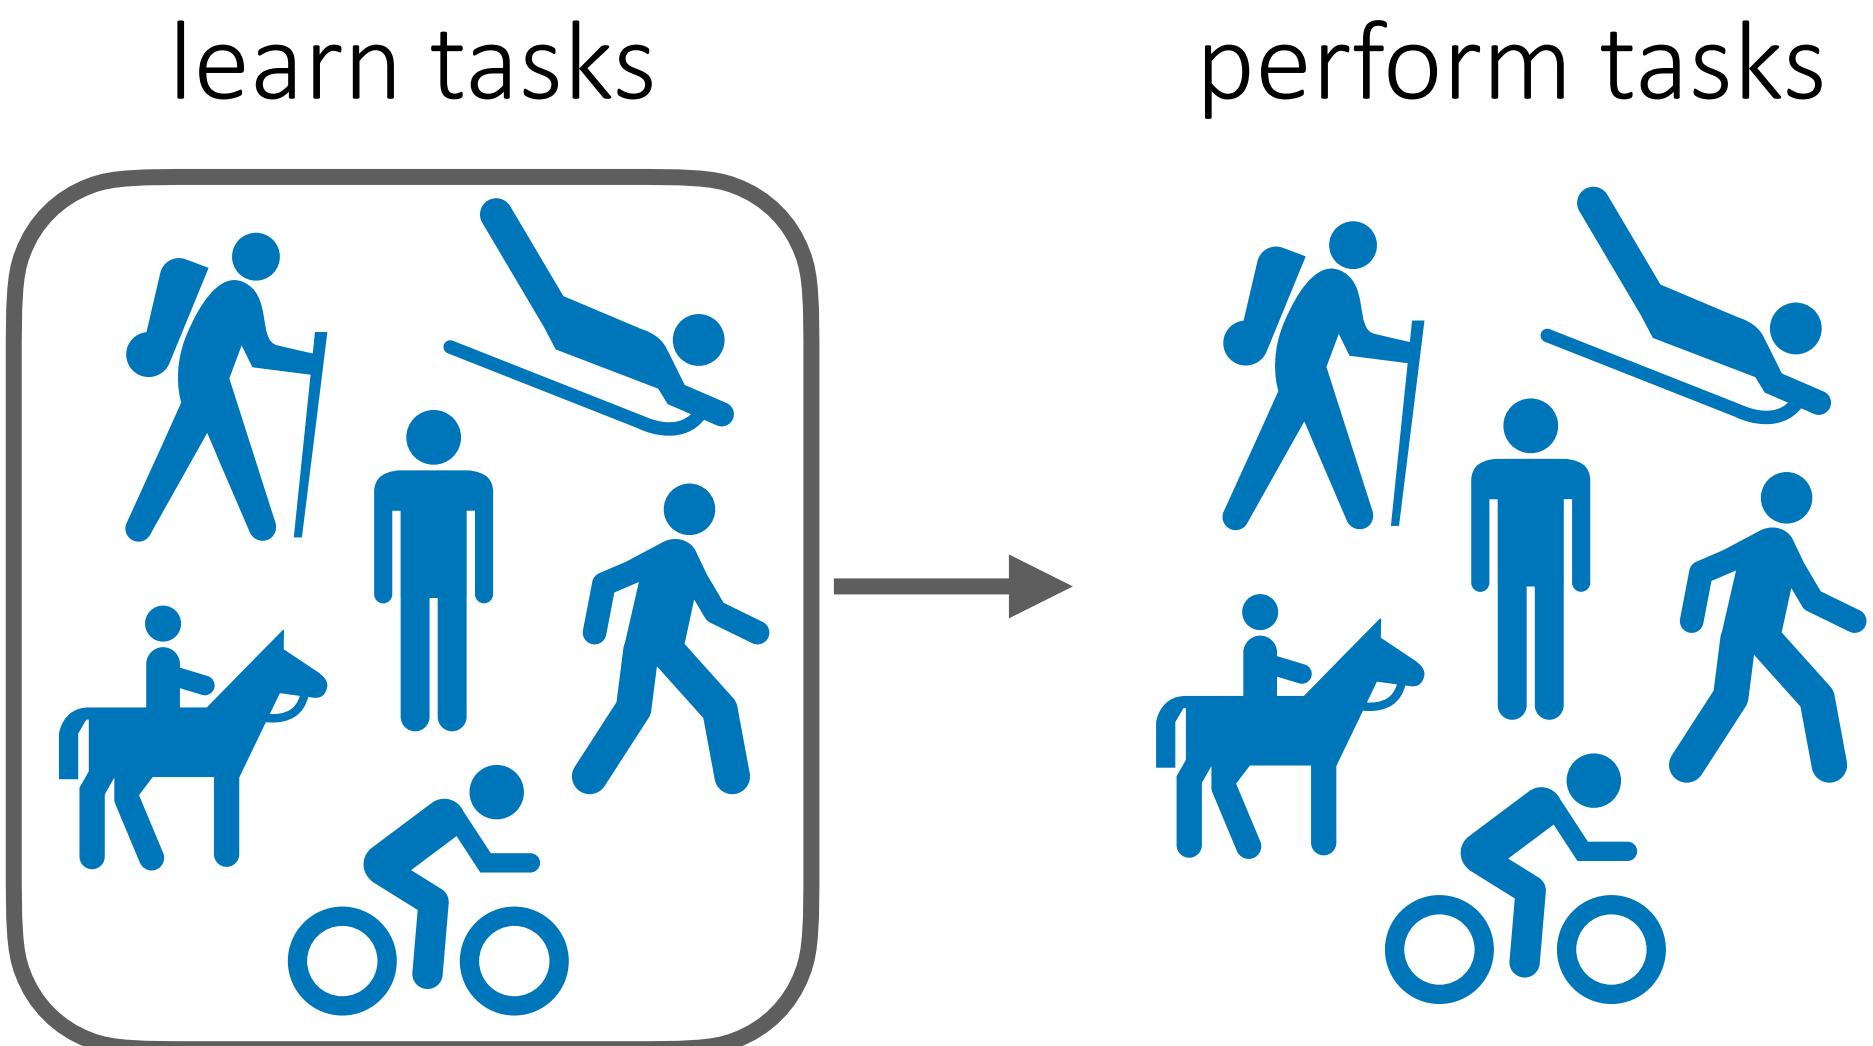
Can we do **better** than the basics?

Revisiting the problem statement
from the **meta-learning perspective**

A brief review of problem statements.

Multi-Task Learning

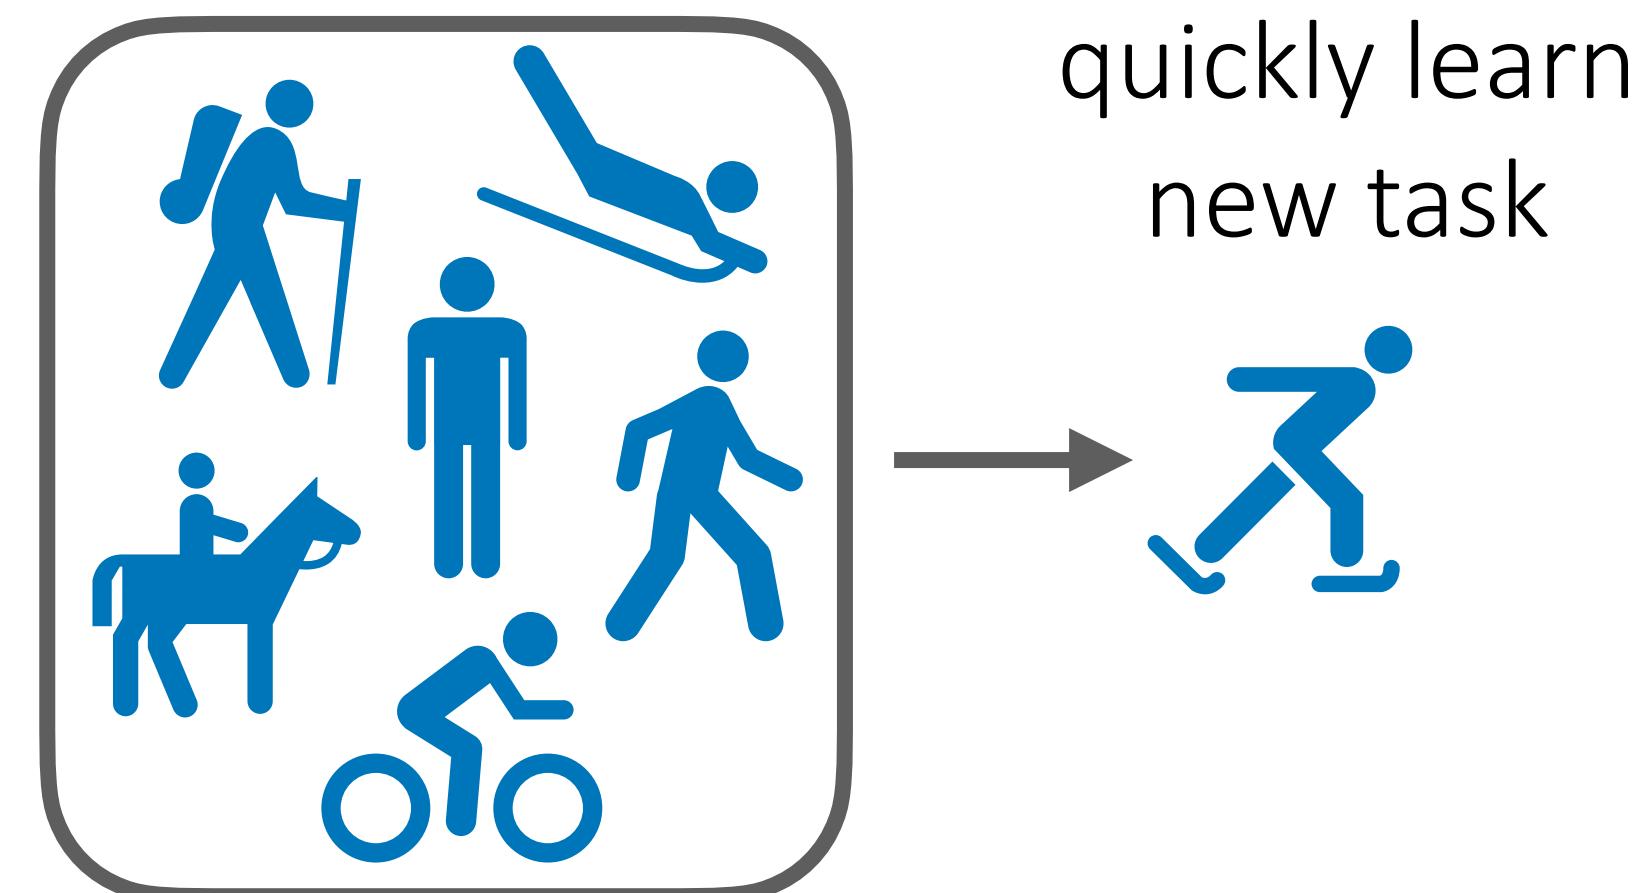
Learn to solve a set of tasks.



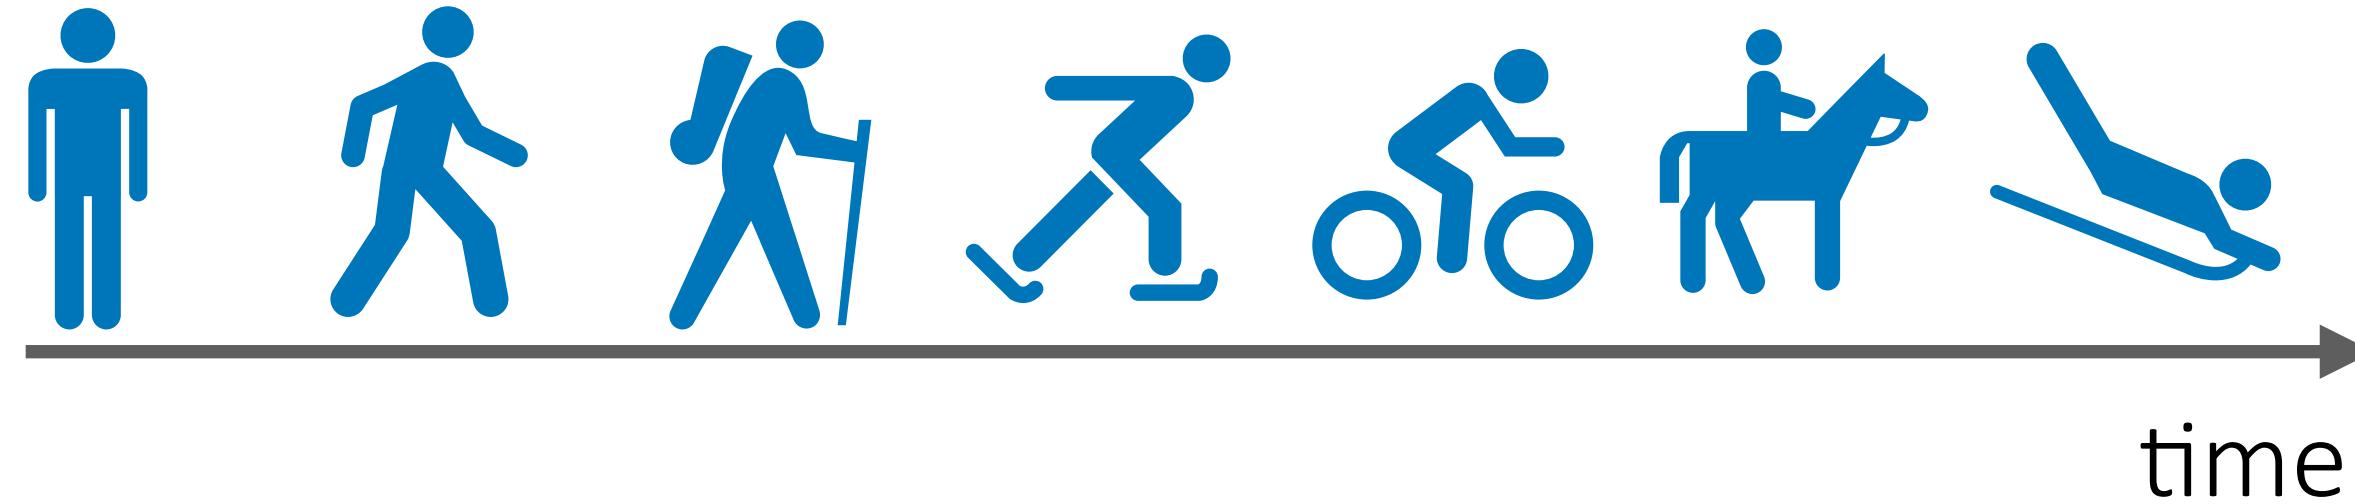
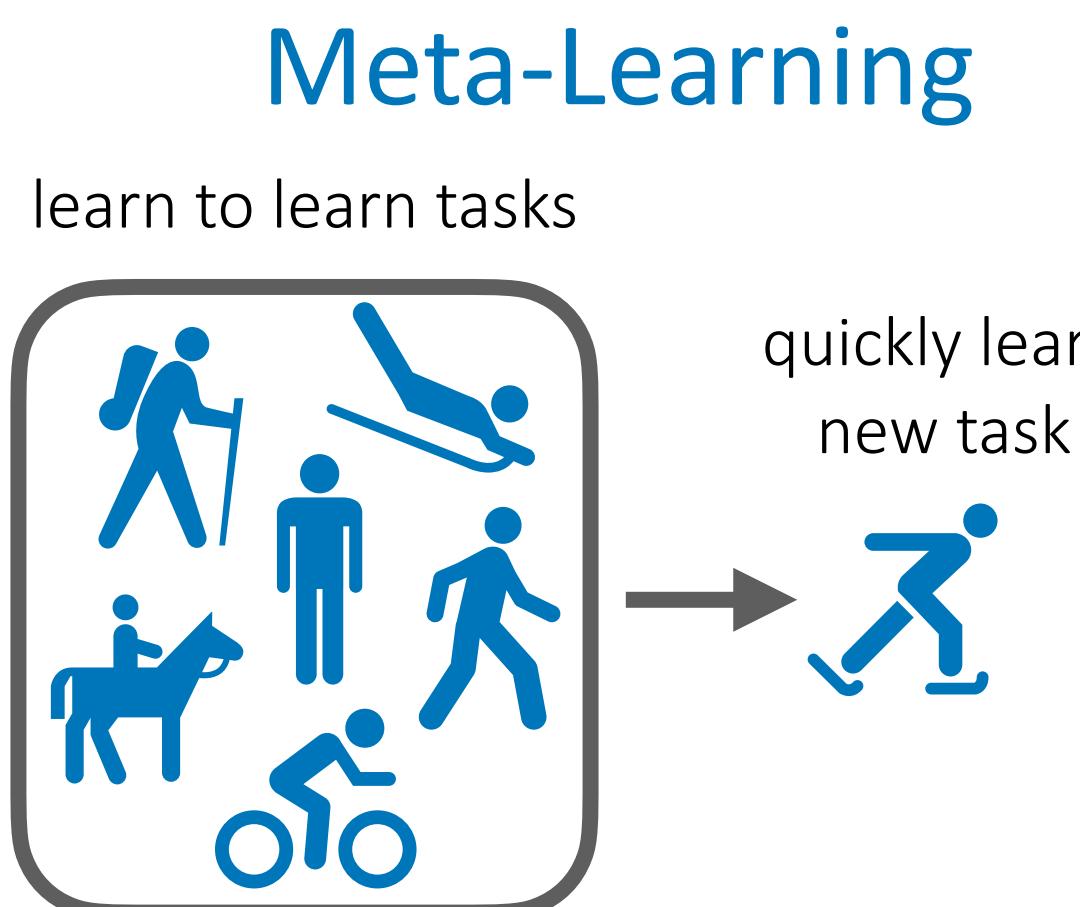
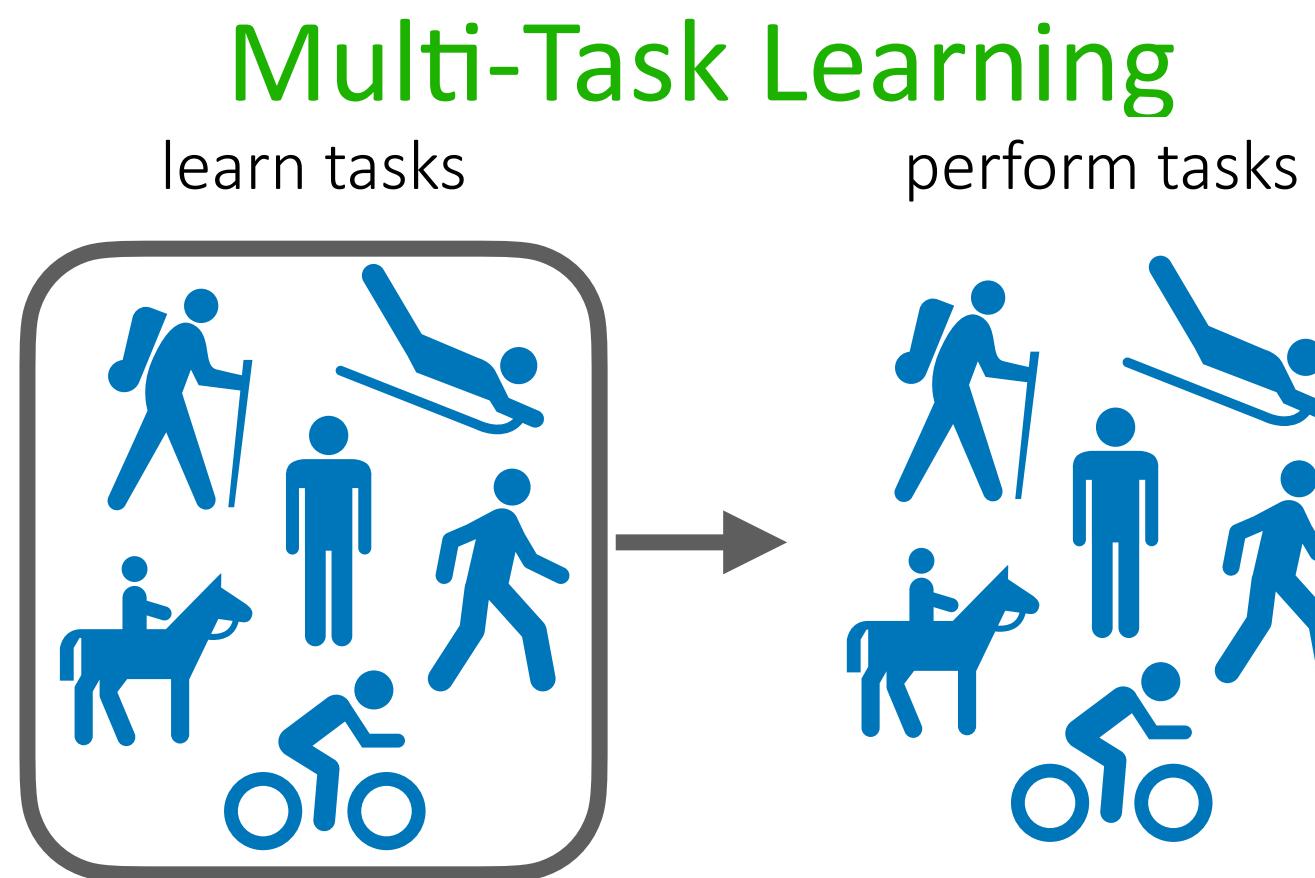
Meta-Learning

Given i.i.d. task distribution,
learn a new task efficiently

learn to learn tasks



In contrast, many real world settings look like:



Our agents may not be given a large batch of data/tasks right off the bat!

Some examples:

- a student learning concepts in school
- a deployed **image classification system** learning from a stream of images from users
- a **robot** acquiring an increasingly large set of skills in different environments
- a **virtual assistant** learning to help different users with different tasks at different points in time
- a **doctor's assistant** aiding in medical decision-making

Some Terminology

Sequential learning settings

online learning, lifelong learning, continual learning, incremental learning, streaming data

distinct from sequence data and sequential decision-making

What is the lifelong learning *problem statement*?

- Exercise:**
1. Pick an example setting.
 2. Discuss problem statement in small groups:
 - (a) how would you set-up an experiment to develop & test your algorithm?
 - (b) what are desirable/required properties of the algorithm?
 - (c) how do you evaluate such a system?

- Example settings:**
- A. a student learning concepts in school
 - B. a deployed image classification system learning from a stream of images from users
 - C. a robot acquiring an increasingly large set of skills in different environments
 - D. a virtual assistant learning to help different users with different tasks at different points in time
 - E. a doctor's assistant aiding in medical decision-making

What is the lifelong learning *problem statement*?

Problem variations:

- task/data order: i.i.d. vs. predictable vs. curriculum vs. adversarial
- discrete task boundaries vs. continuous shifts (vs. both)
- known task boundaries/shifts vs. unknown

Some considerations:

- model performance
- data efficiency
- computational resources
- memory
- others: privacy, interpretability, fairness, test time compute & memory

Substantial variety in problem statement!

What is the lifelong learning *problem statement*?

General [supervised] online learning problem:

for $t = 1, \dots, n$

observe x_t \leftarrow if **observable task boundaries**: observe x_t, z_t

predict \hat{y}_t

observe label y_t

i.i.d. setting: $x_t \sim p(x)$, $y_t \sim p(y|x)$

p not a function of t

otherwise: $x_t \sim p_t(x)$, $y_t \sim p_t(y|x)$

streaming setting: cannot store (x_t, y_t)

- lack of memory
- lack of computational resources
- privacy considerations
- want to study neural memory mechanisms

true in some cases, but not in many cases!

- 9 - recall: replay buffers

What do you want from your lifelong learning algorithm?

minimal regret (that grows slowly with t)

regret: cumulative loss of learner — cumulative loss of best learner in hindsight

$$\text{Regret}_T := \sum_1^T \mathcal{L}_t(\theta_t) - \min_{\theta} \sum_1^T \mathcal{L}_t(\theta)$$

(cannot be evaluated in practice, useful for analysis)

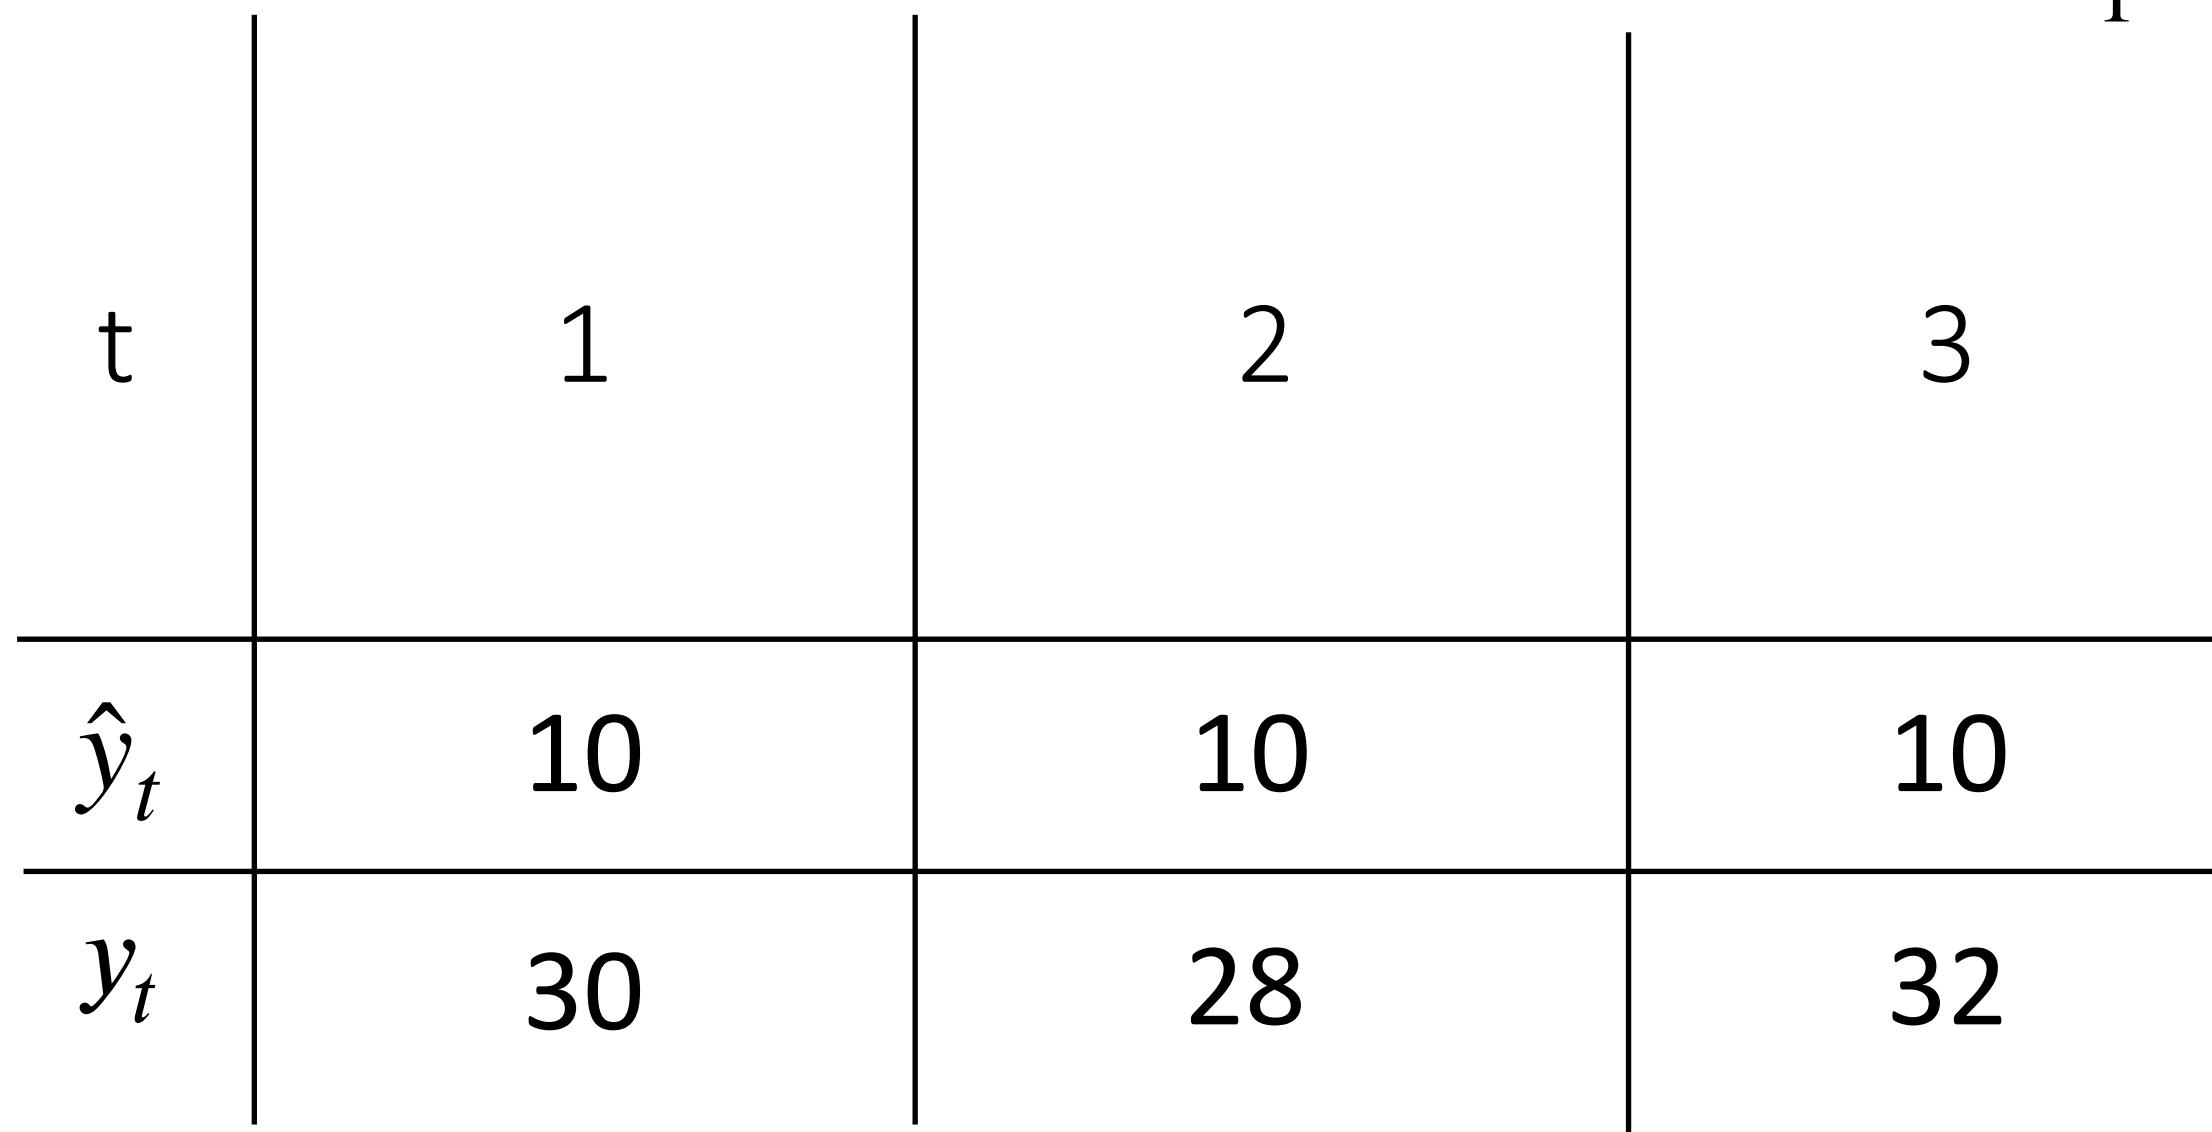
Regret that grows linearly in t is trivial. Why?

What do you want from your lifelong learning algorithm?

minimal regret (that grows slowly with t)

regret: cumulative loss of learner — cumulative loss of best learner in hindsight

$$\text{Regret}_T := \sum_1^T \mathcal{L}_t(\theta_t) - \min_{\theta} \sum_1^T \mathcal{L}_t(\theta)$$

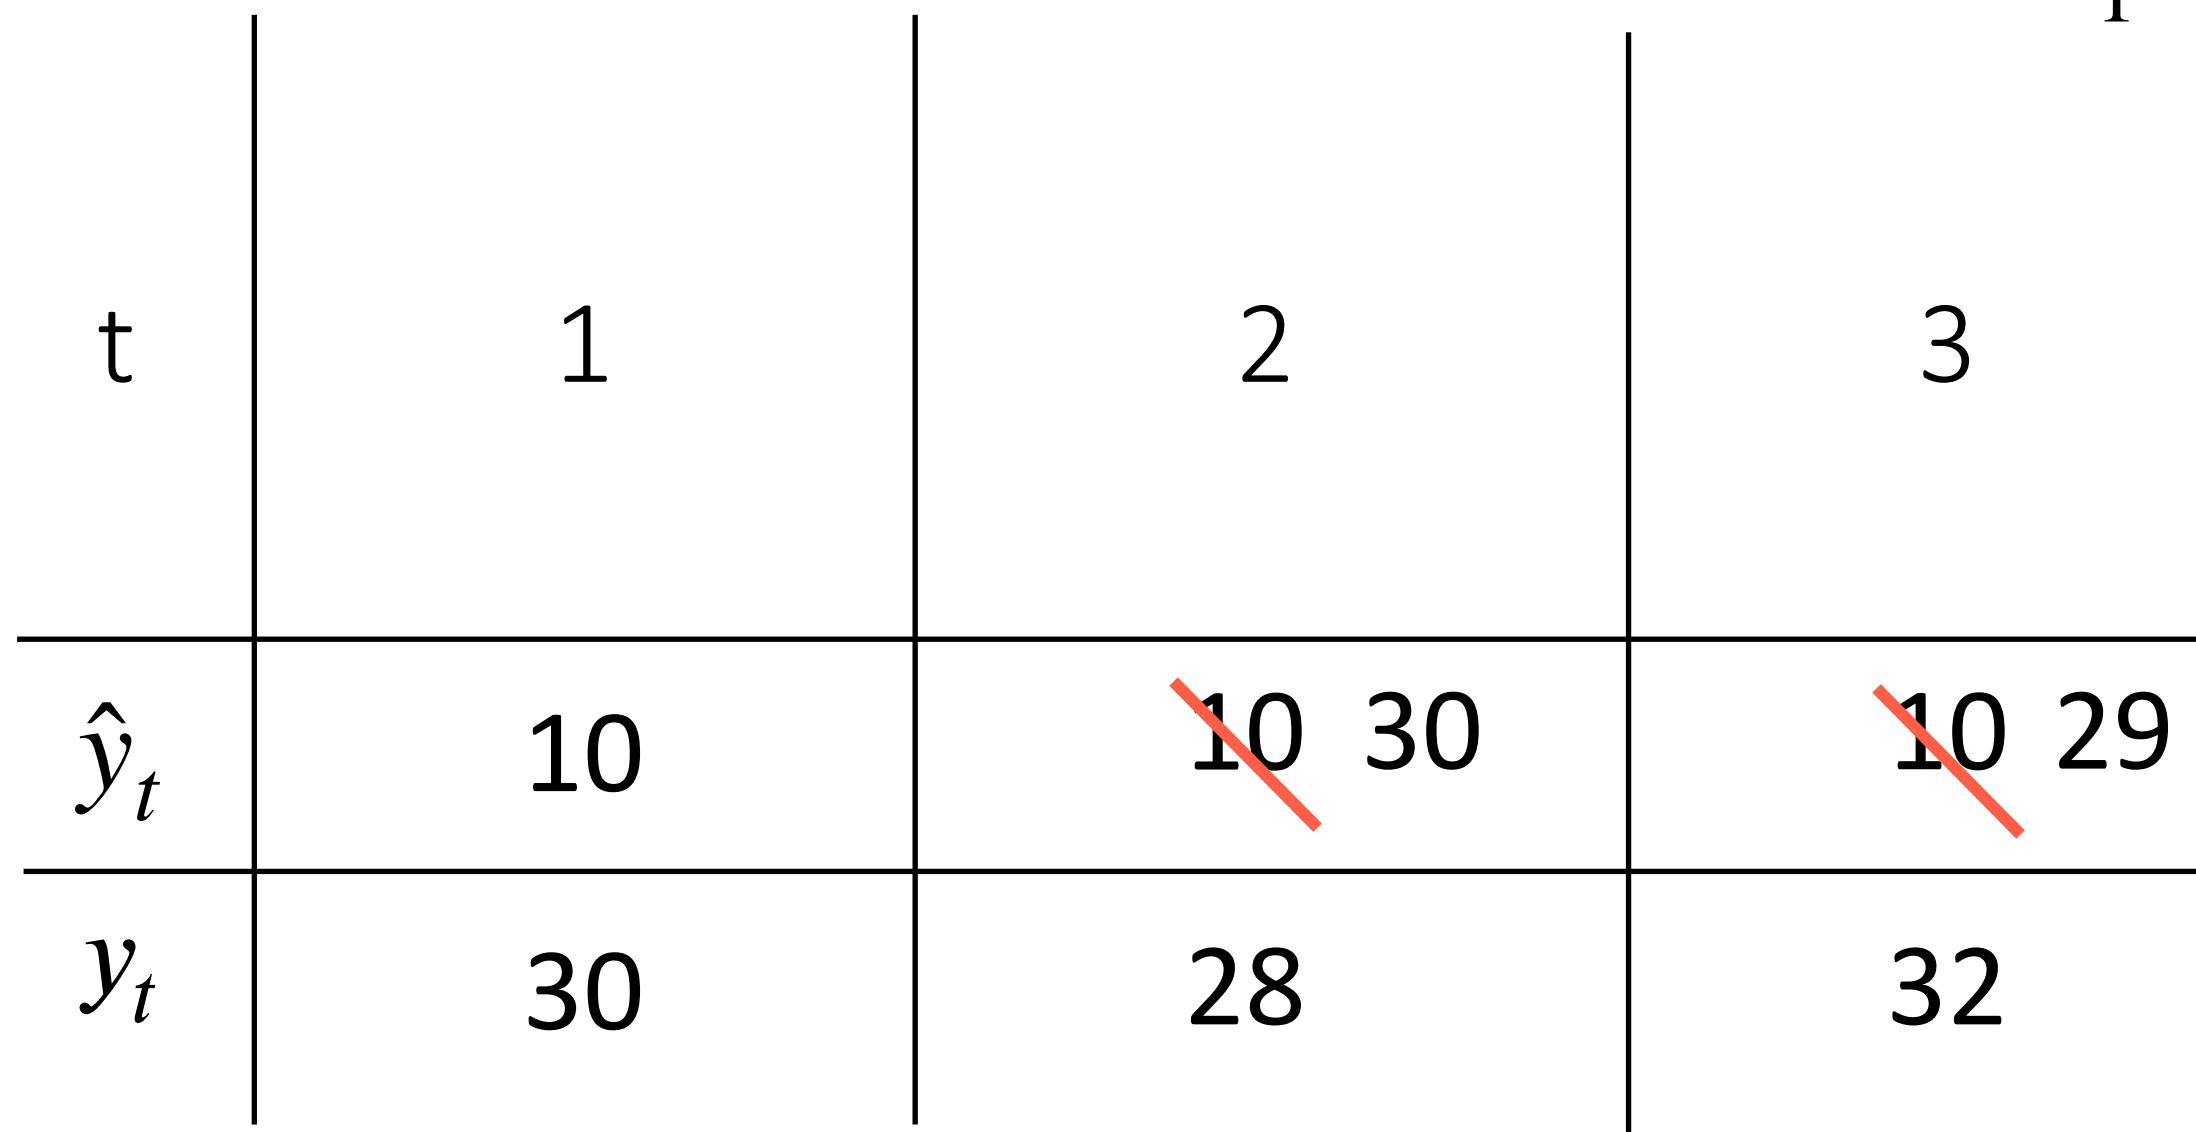


What do you want from your lifelong learning algorithm?

minimal regret (that grows slowly with t)

regret: cumulative loss of learner — cumulative loss of best learner in hindsight

$$\text{Regret}_T := \sum_1^T \mathcal{L}_t(\theta_t) - \min_{\theta} \sum_1^T \mathcal{L}_t(\theta)$$



What do you want from your lifelong learning algorithm?

positive & negative transfer

positive **forward** transfer: previous tasks cause you to do better on future tasks
compared to learning future tasks from scratch

positive **backward** transfer: current tasks cause you to do better on previous tasks
compared to learning past tasks from scratch

positive -> negative : better -> worse

Plan for Today

The lifelong learning problem statement

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from the meta-learning perspective

Approaches

Store all the data you've seen so far, and train on it. → follow the leader algorithm

- + will achieve very strong performance
- computation intensive → Continuous fine-tuning can help.
- can be memory intensive [depends on the application]

Take a gradient step on the datapoint you observe. → stochastic gradient descent

- + computationally cheap
- + requires 0 memory
- subject to negative backward transfer
“forgetting”
sometimes referred to as
catastrophic forgetting
- slow learning

Can we do better?

Applying a simple continual learning algorithm to robotics



86%

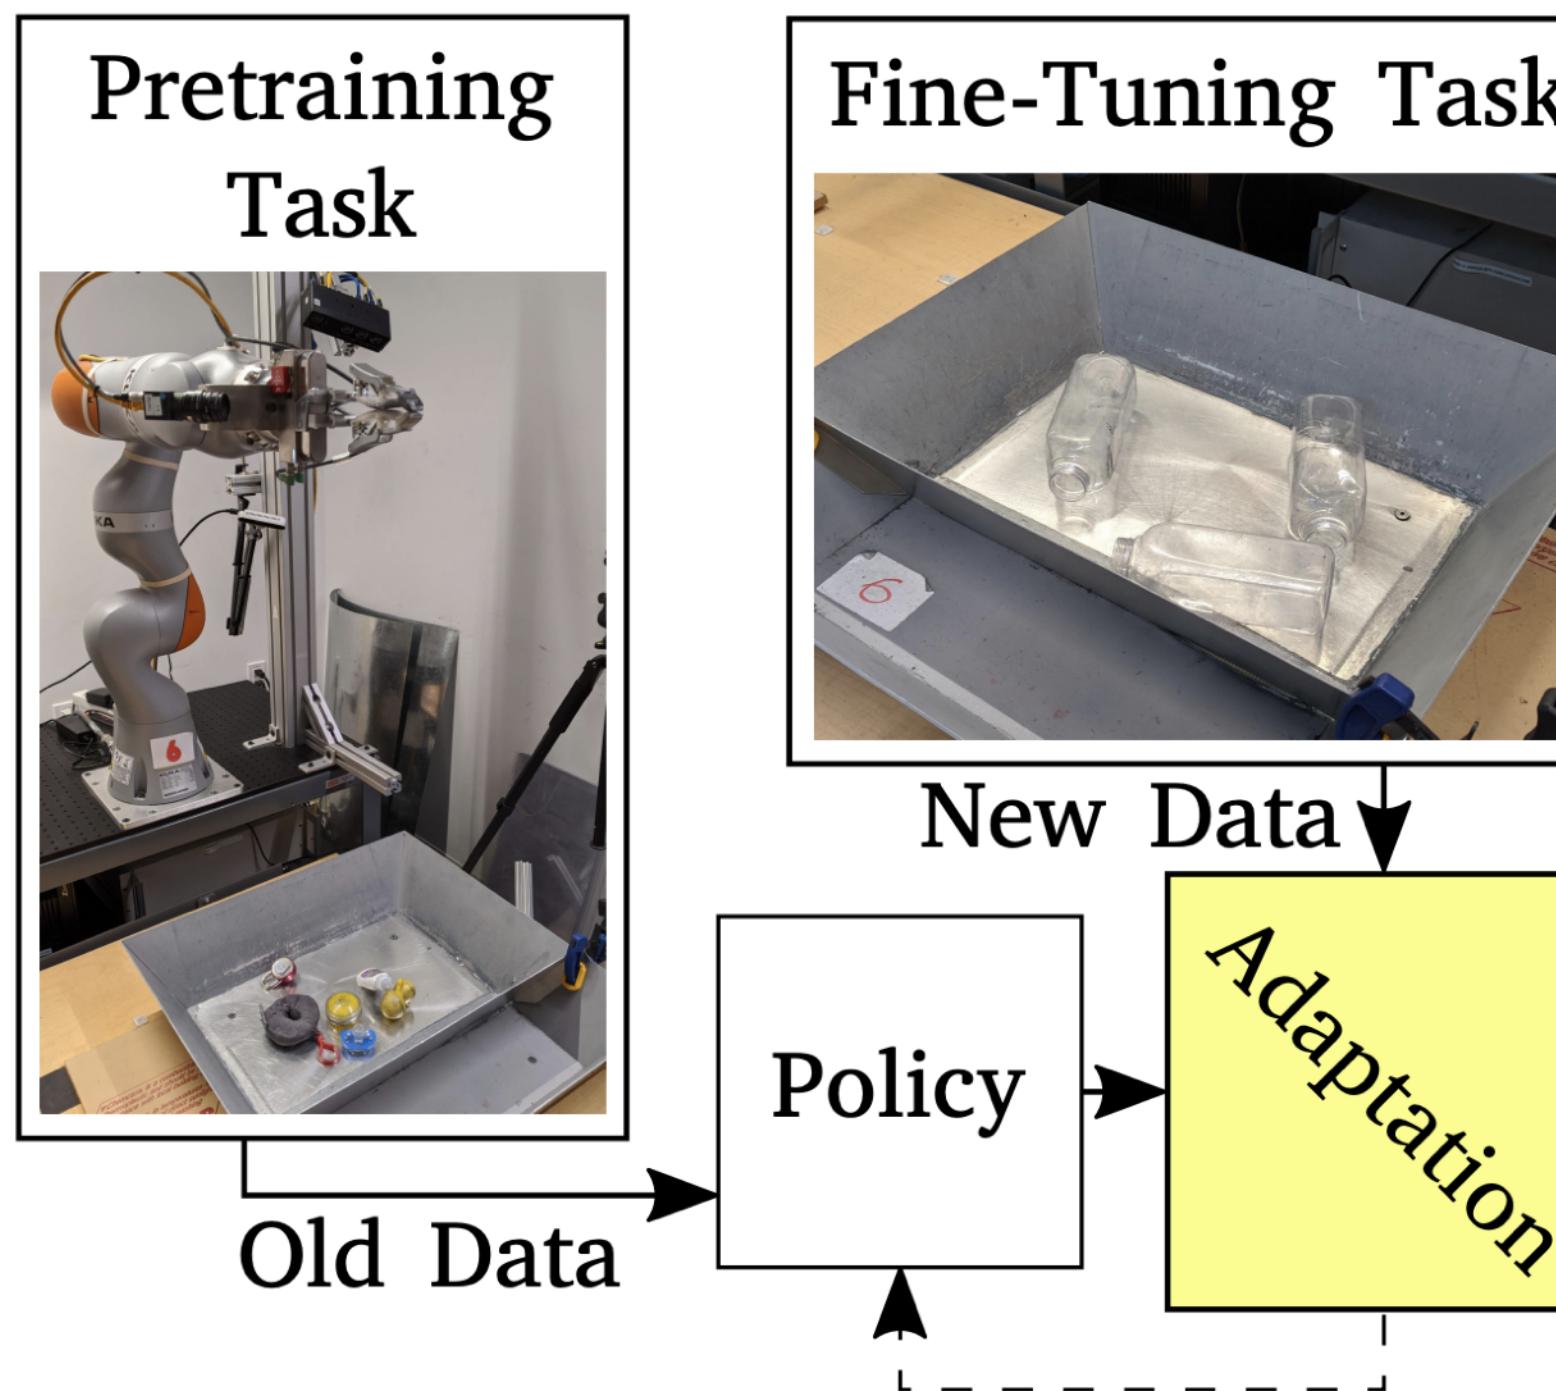


49%

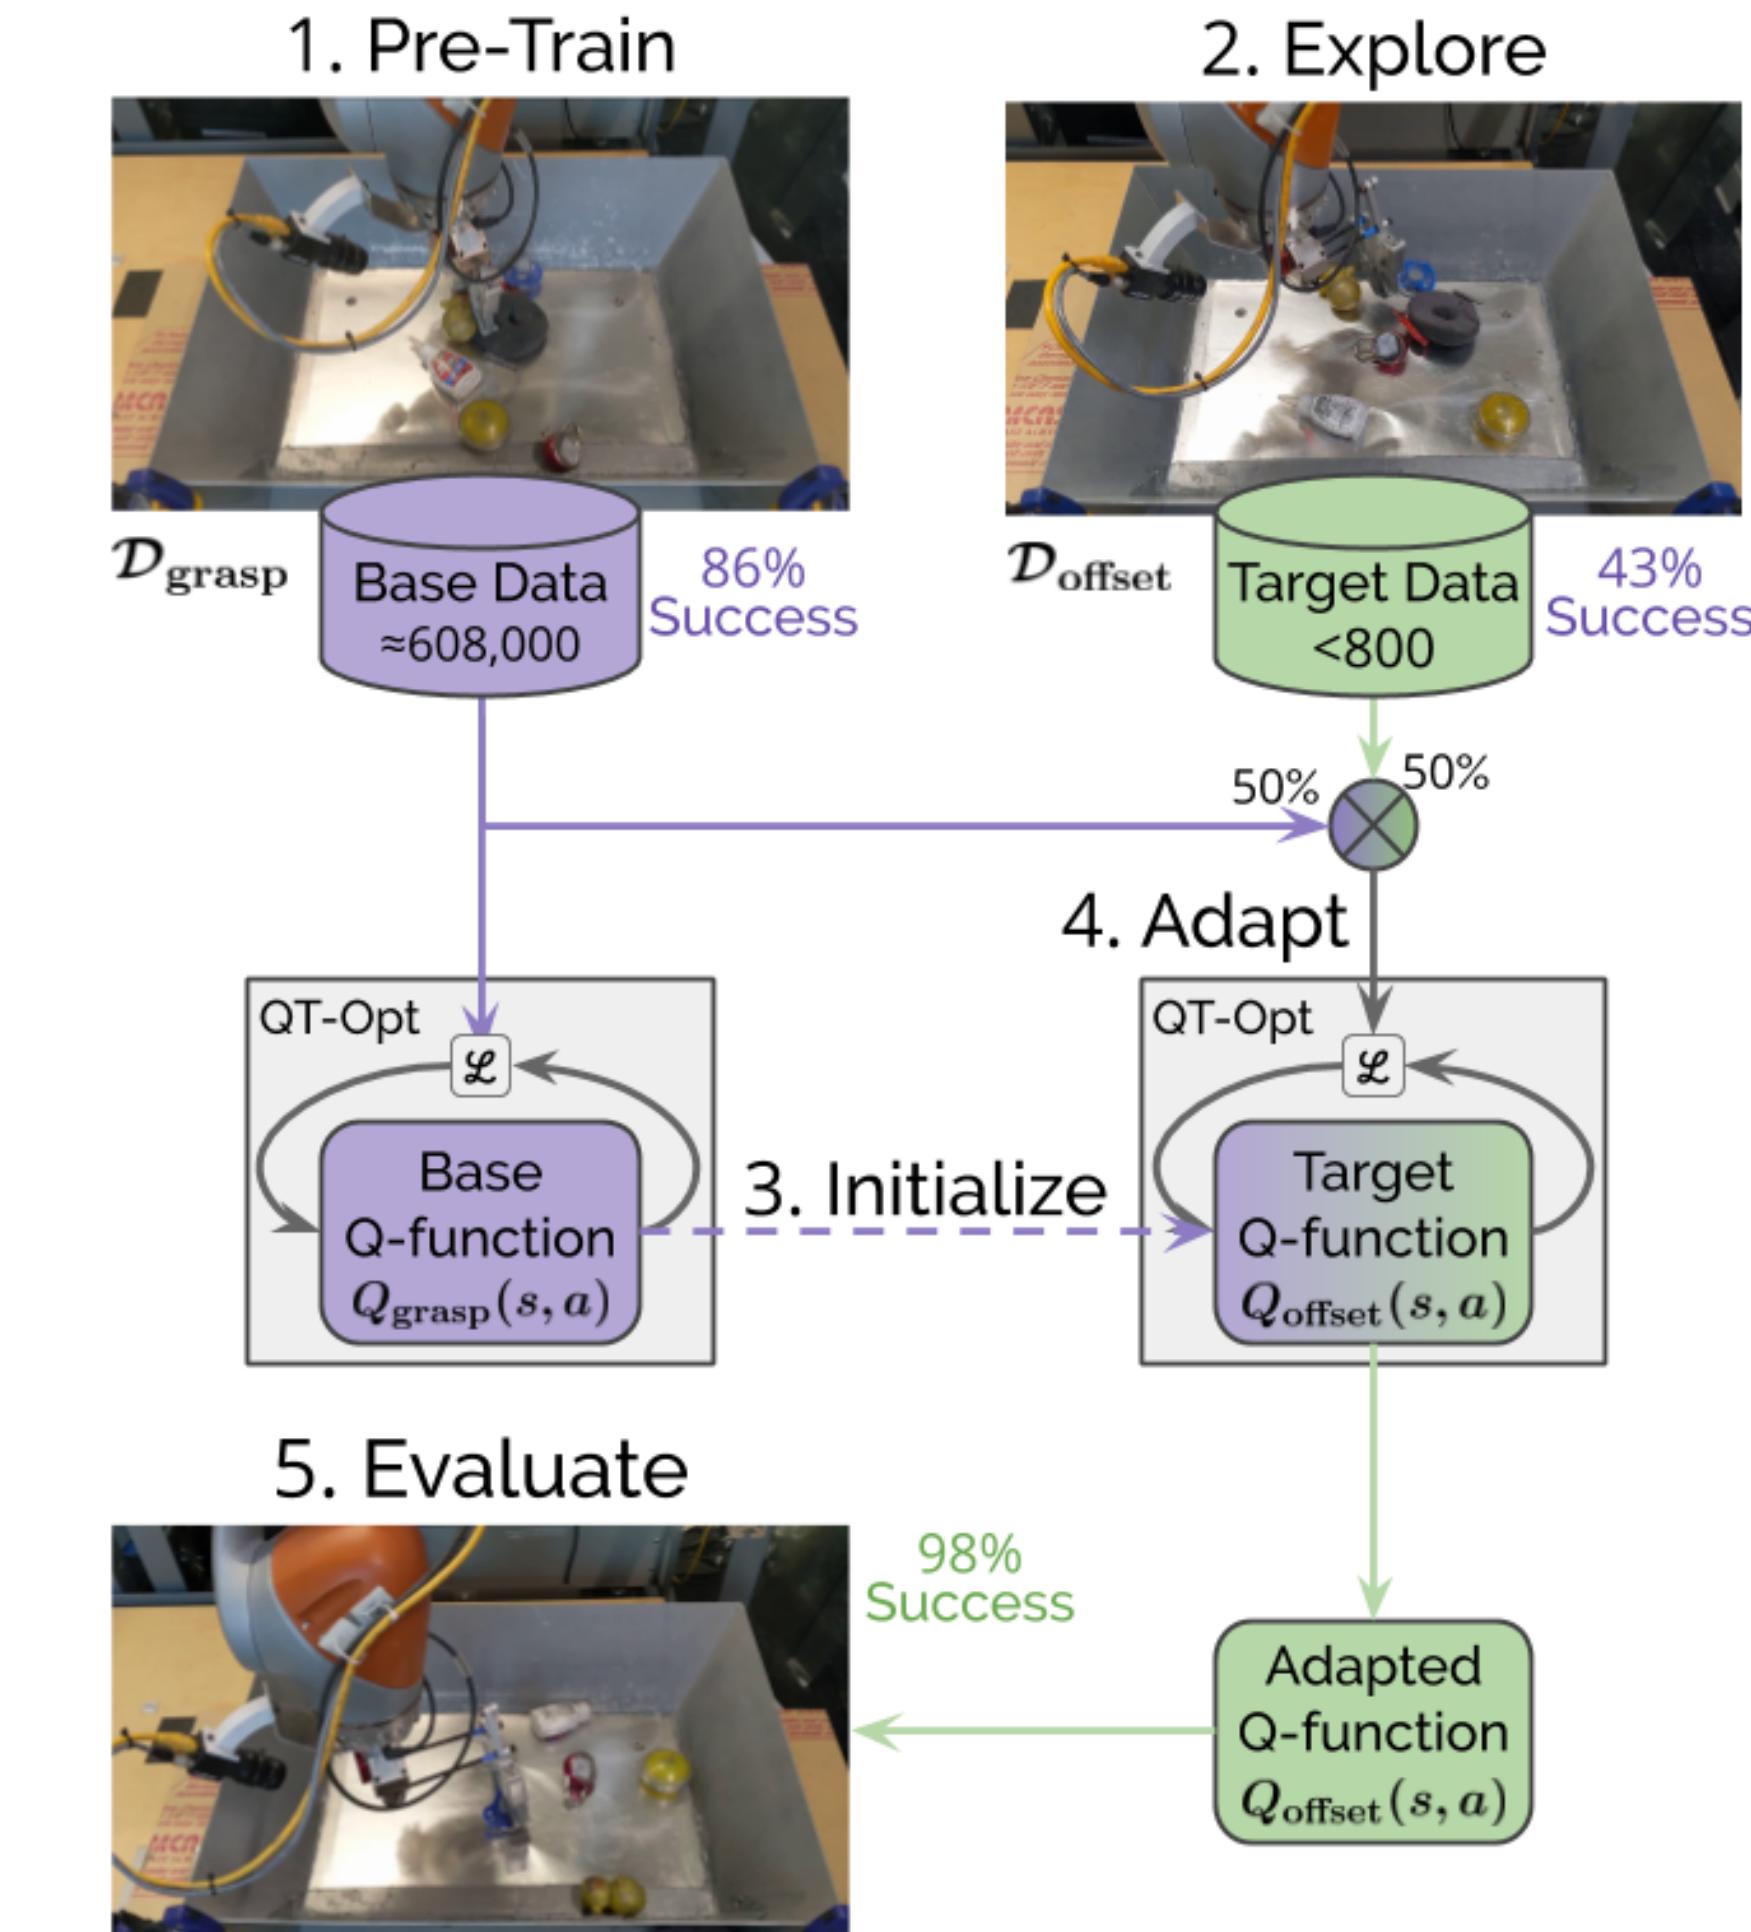
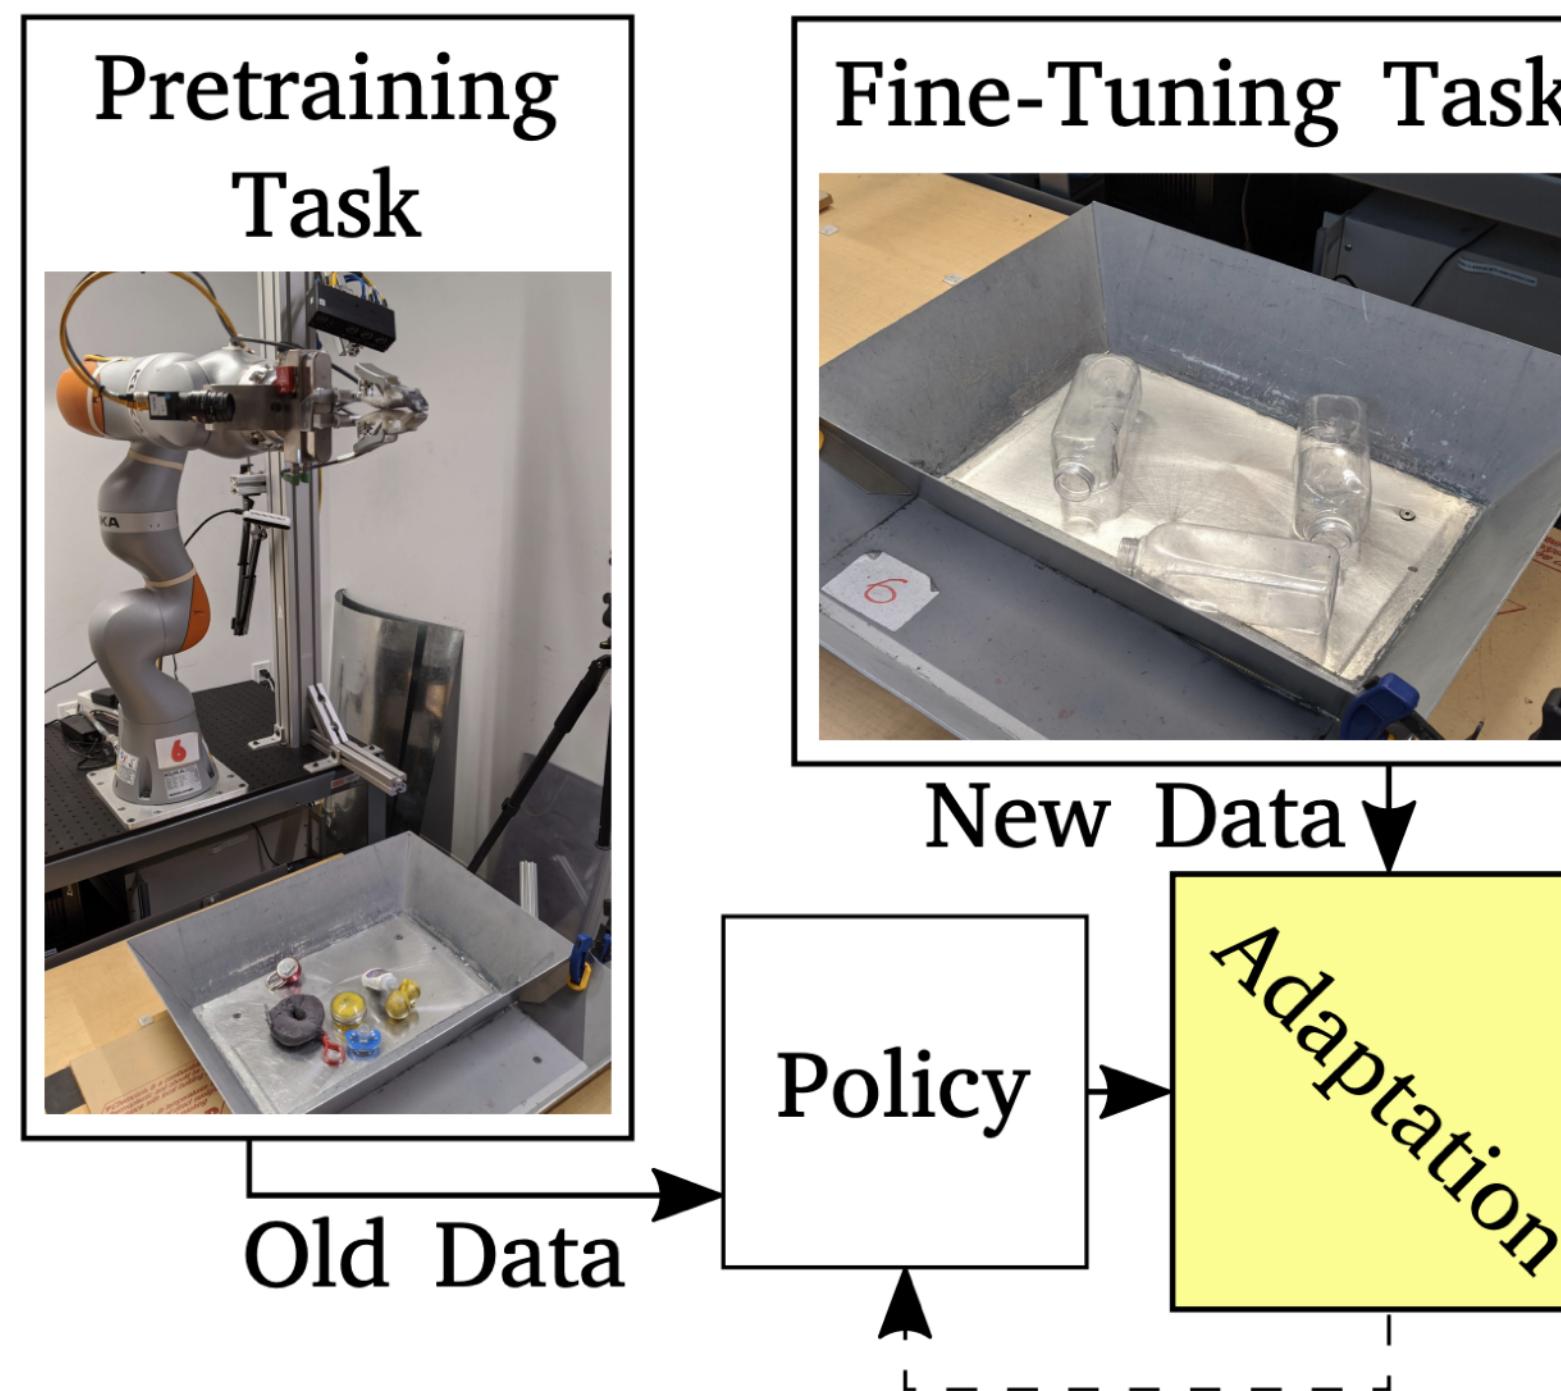


7 robots collected 580k grasps

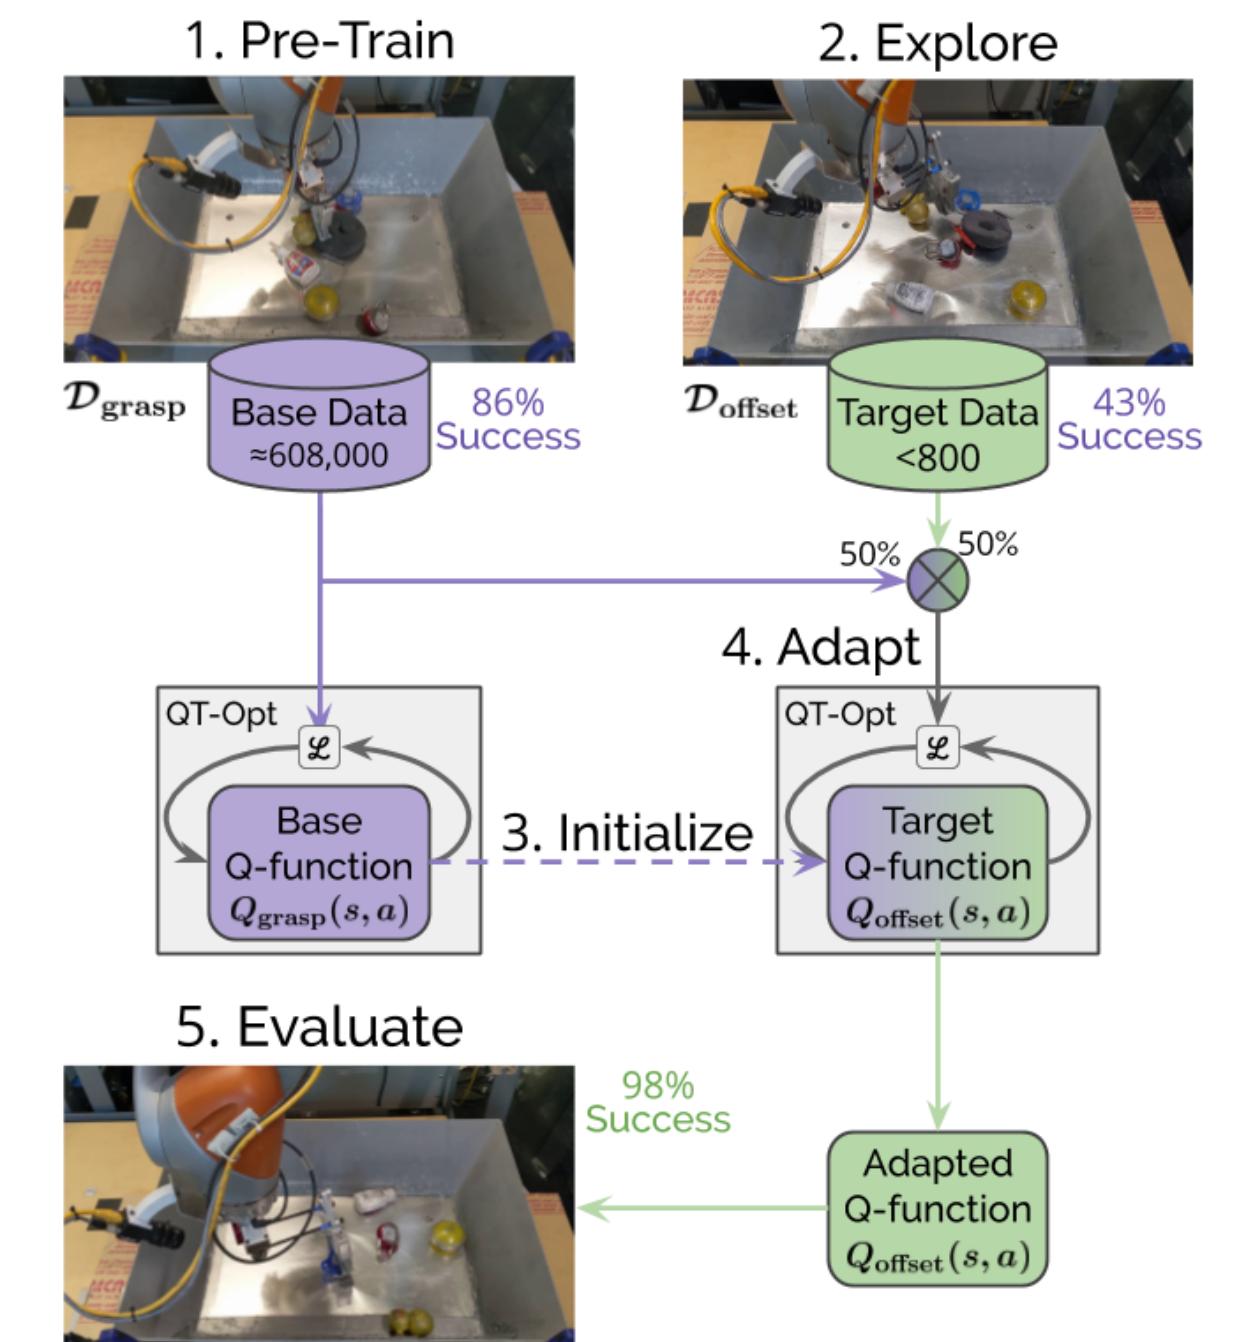
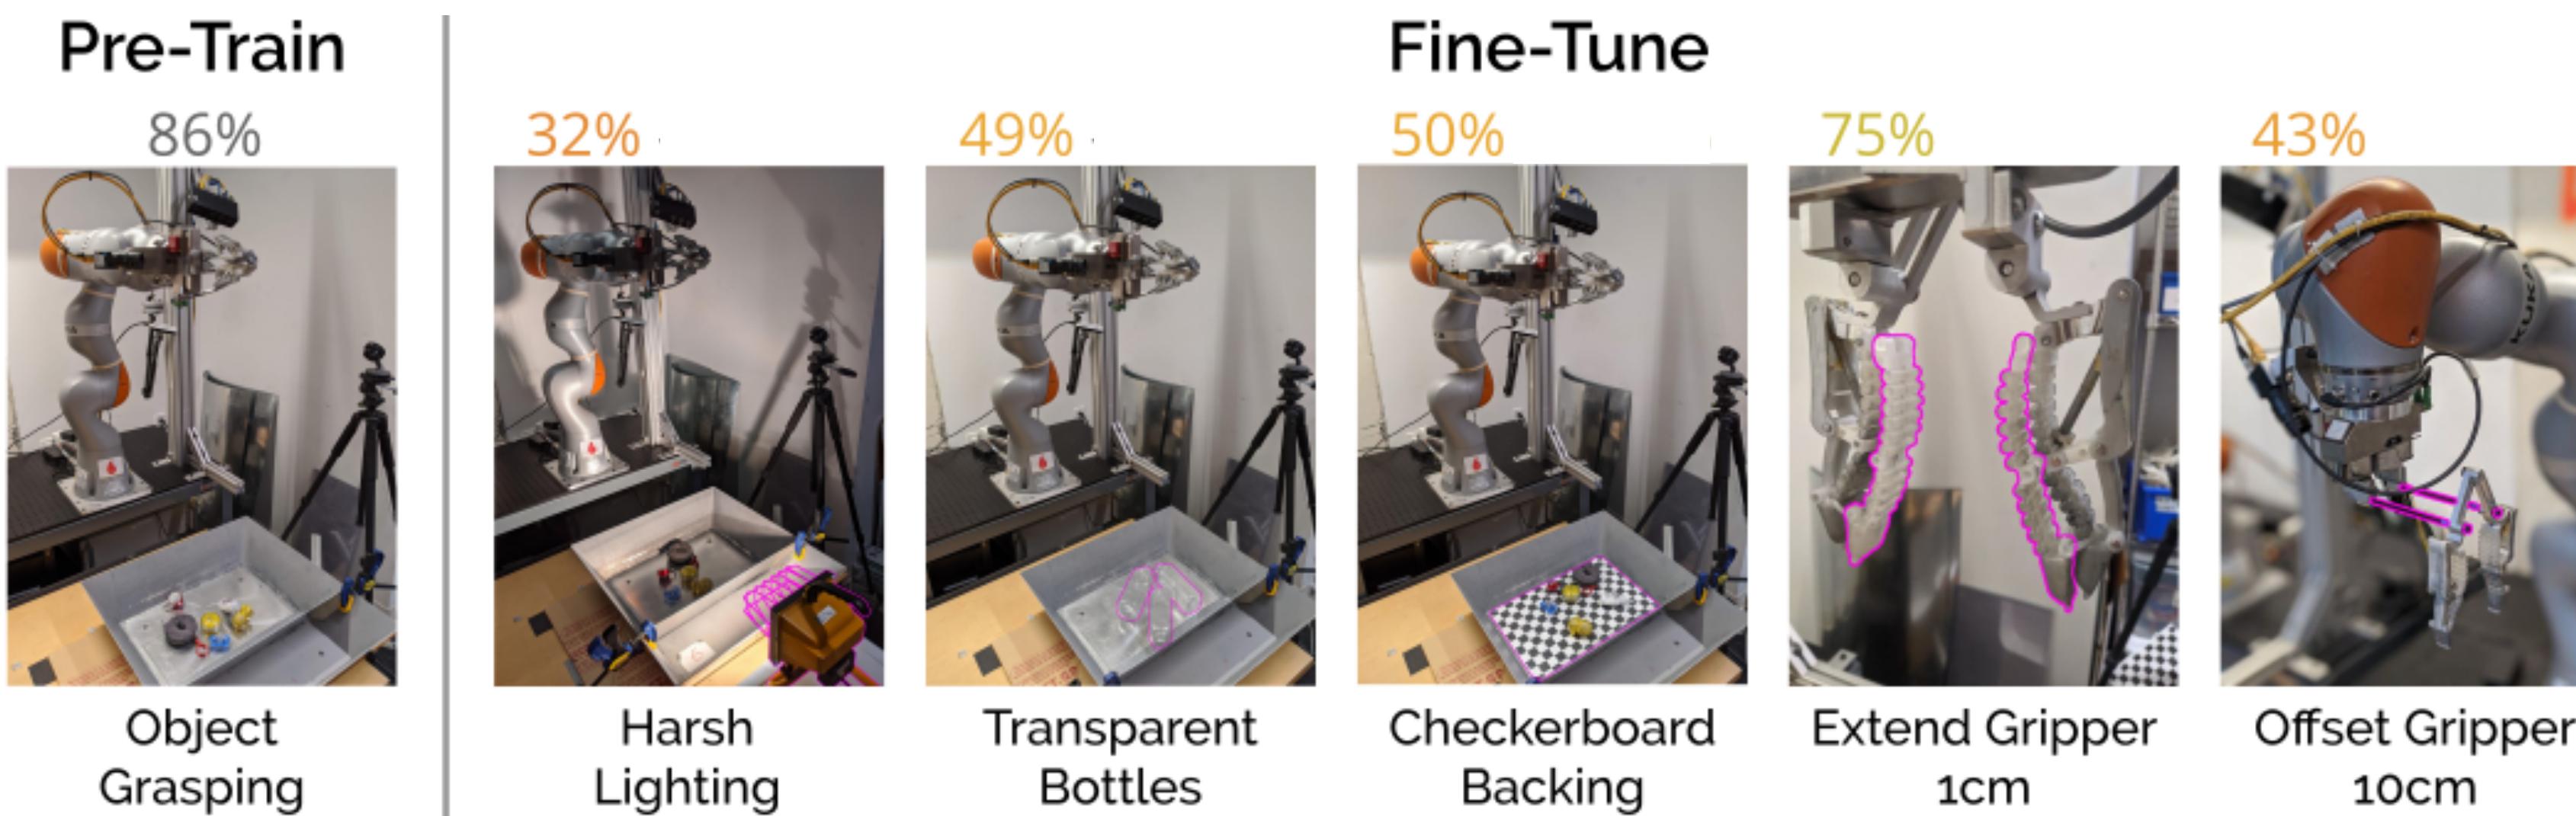
Applying a simple continual learning algorithm to robotics



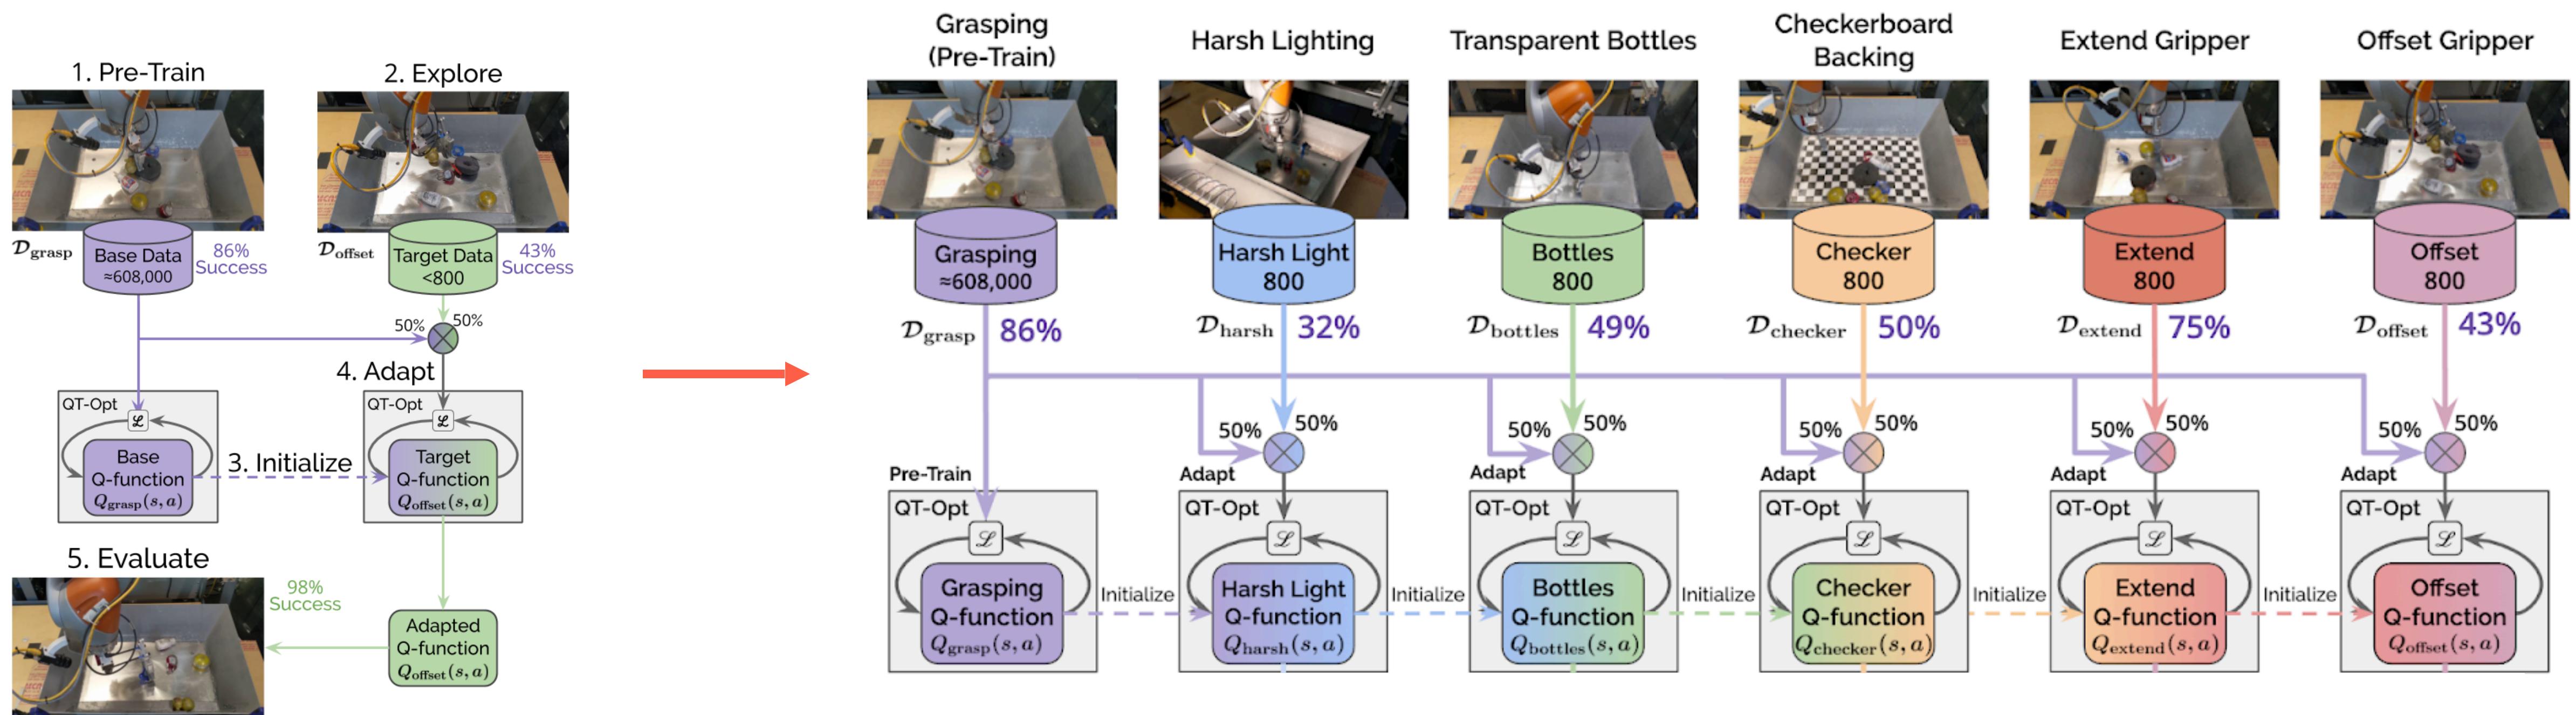
Applying a simple continual learning algorithm to robotics



Applying a simple continual learning algorithm to robotics



Applying a simple continual learning algorithm to robotics



What about backward transfer?

Can we do better?

Plan for Today

The lifelong learning problem statement

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from the meta-learning perspective

Case Study: Can we modify vanilla SGD to avoid negative backward transfer?
(from scratch)

Idea:

- (1) store small amount of data per task in memory
- (2) when making updates for new tasks, ensure that they don't **unlearn** previous tasks

How do we accomplish (2)?

learning predictor $y_t = f_\theta(x_t, z_t)$ memory: \mathcal{M}_k for task z_k

For $t = 0, \dots, T$

minimize $\mathcal{L}(f_\theta(\cdot, z_t), (x_t, y_t))$

subject to $\mathcal{L}(f_\theta, \mathcal{M}_k) \leq \mathcal{L}(f_\theta^{t-1}, \mathcal{M}_k)$ for all $k < t$

(i.e. s.t. loss on previous tasks doesn't get worse)

Assume local linearity:

$$\langle g_t, g_k \rangle := \left\langle \frac{\partial \mathcal{L}(f_\theta, (x_t, y_t))}{\partial \theta}, \frac{\partial \mathcal{L}(f_\theta, \mathcal{M}_k)}{\partial \theta} \right\rangle \geq 0 \quad \text{for all } z_k < z_t$$

Can formulate & solve as a QP.

Experiments

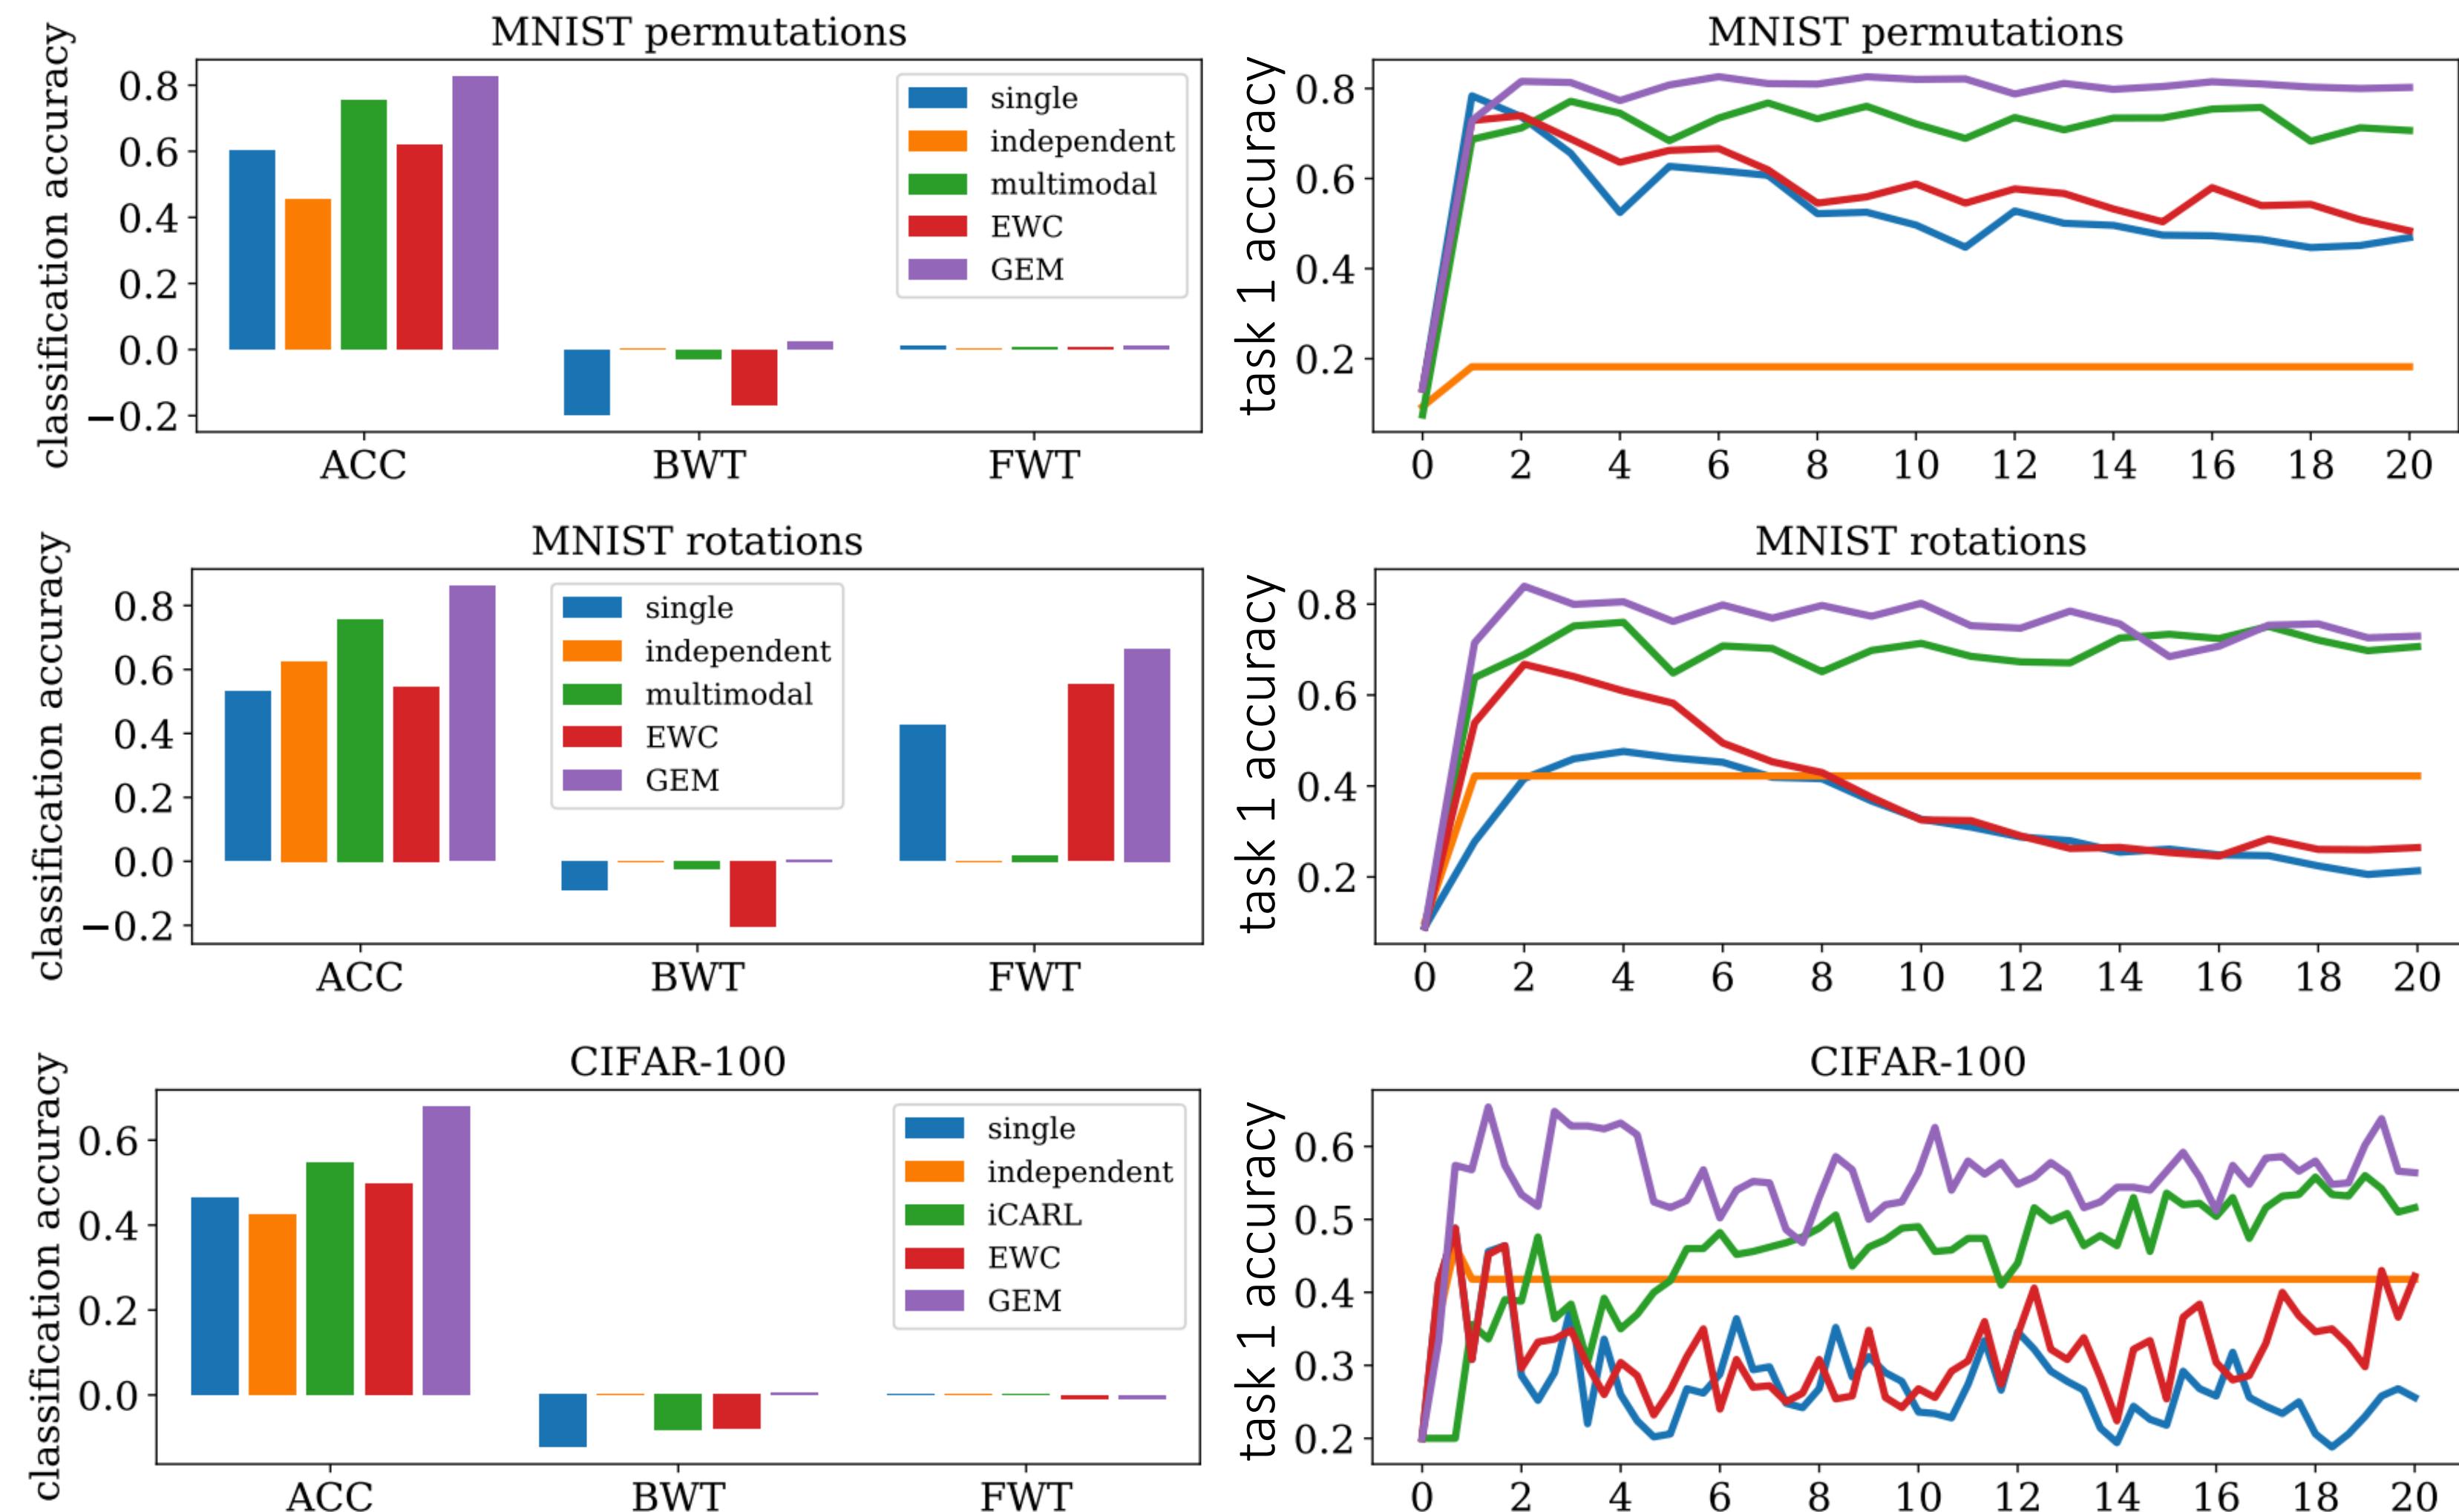
Problems:

- MNIST permutations
- MNIST rotations
- CIFAR-100 (5 new classes/task)

BWT: backward transfer,

FWT: forward transfer

Total memory size:
5012 examples



If we take a step back... do these experimental domains make sense?

Can we meta-learn how to avoid negative backward transfer?

Javed & White. *Meta-Learning Representations for Continual Learning*. NeurIPS '19
Beaulieu et al. *Learning to Continually Learn*. '20

Plan for Today

The lifelong learning problem statement

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from **the meta-learning perspective**

Formulation of online learning when faced with sequence of tasks

Online Learning

(Hannan '57, Zinkevich '03)

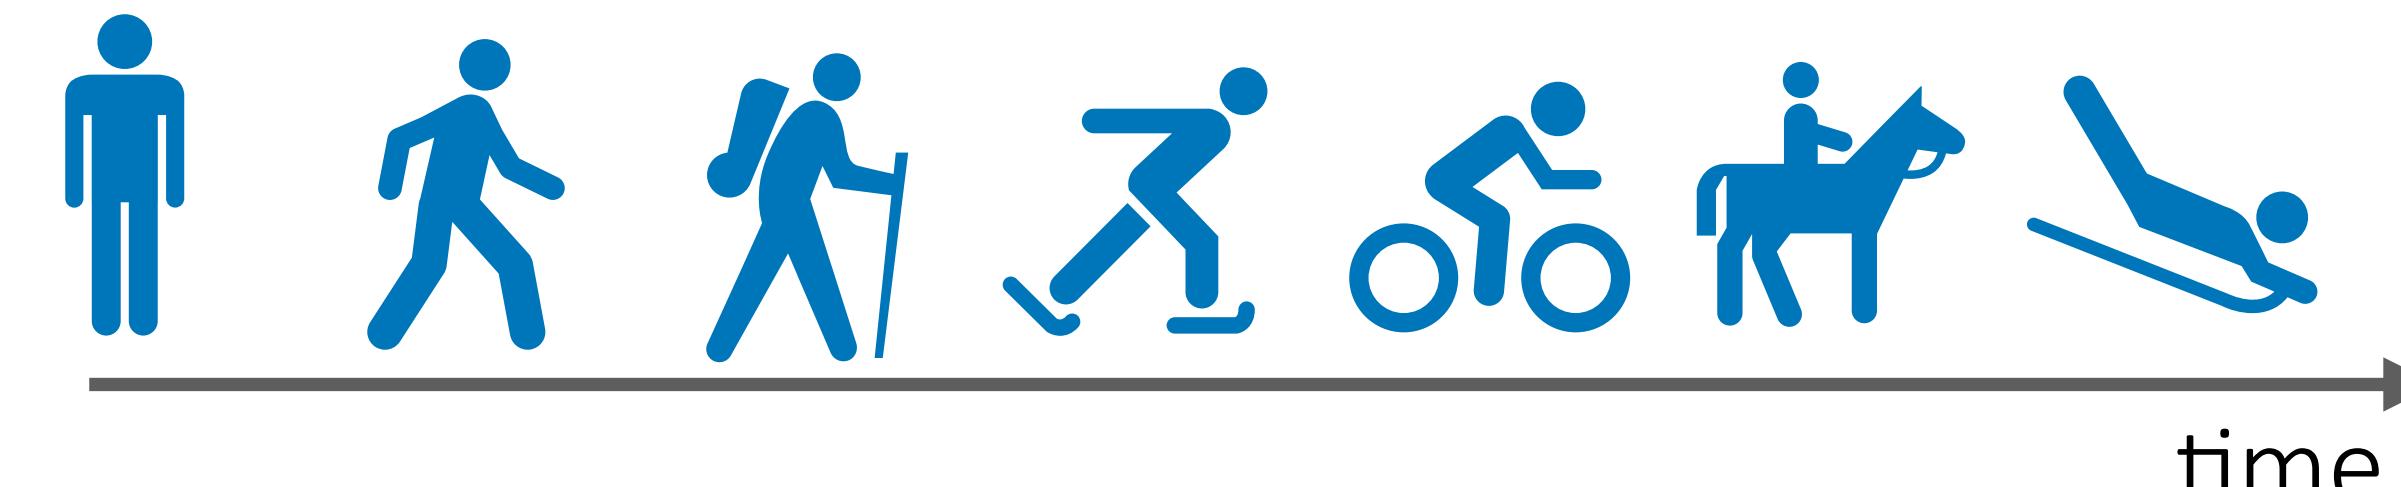
Perform sequence of tasks
while minimizing static regret.

perform perform perform perform perform perform perform



More realistically:

learn learn learn learn learn learn learn



slow learning —————→ rapid learning

Formulation of online learning when faced with sequence of tasks

Online Learning

(Hannan '57, Zinkevich '03)

Perform sequence of tasks
while minimizing static regret.

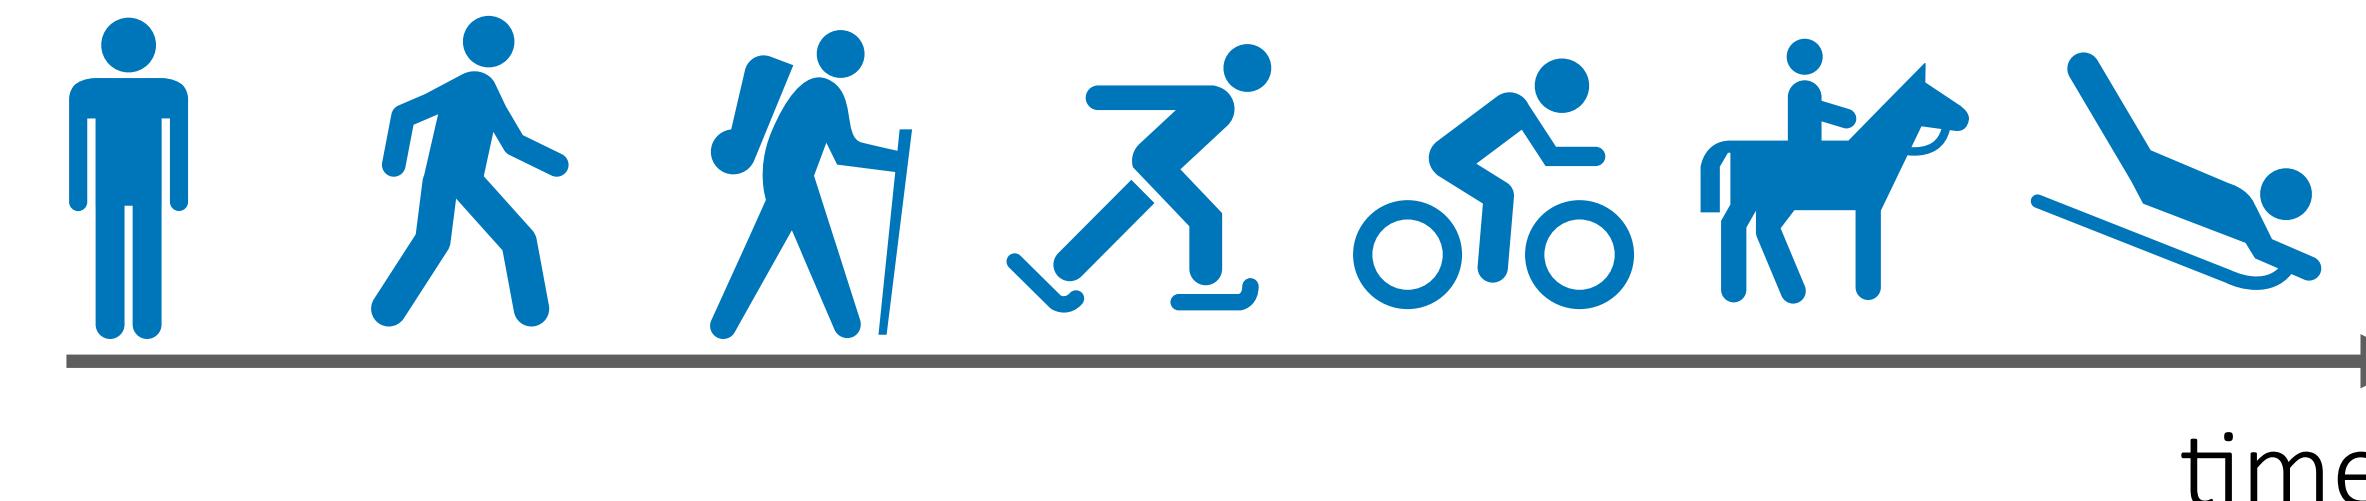
perform perform perform perform perform perform perform



Online Meta-Learning

Efficiently learn a sequence of tasks
from a non-stationary distribution.

learn learn learn learn learn learn learn



evaluate performance after seeing a small amount of data

Primarily a difference in *evaluation*, rather than the *data stream*.

The Online Meta-Learning Setting

for task $t = 1, \dots, n$

observe $\mathcal{D}_t^{\text{tr}}$

use update procedure $\Phi(\theta_t, \mathcal{D}_t^{\text{tr}})$ to produce parameters ϕ_t

observe x_t

predict $\hat{y}_t = f_{\phi_t}(x_t)$

Standard online learning setting

observe label y_t

Goal: Learning algorithm with sub-linear

Loss of algorithm

$$\text{Regret}_T := \sum_{t=1}^T \ell_t(\Phi_t(\theta_t)) - \min_{\theta \in \Theta} \sum_{t=1}^T \ell_t(\Phi_t(\theta))$$

Loss of best algorithm
in hindsight

Can we apply meta-learning in lifelong learning settings?

Recall the **follow the leader** (FTL) algorithm:

Store all the data you've seen so far, and train on it.

Deploy model on current task.

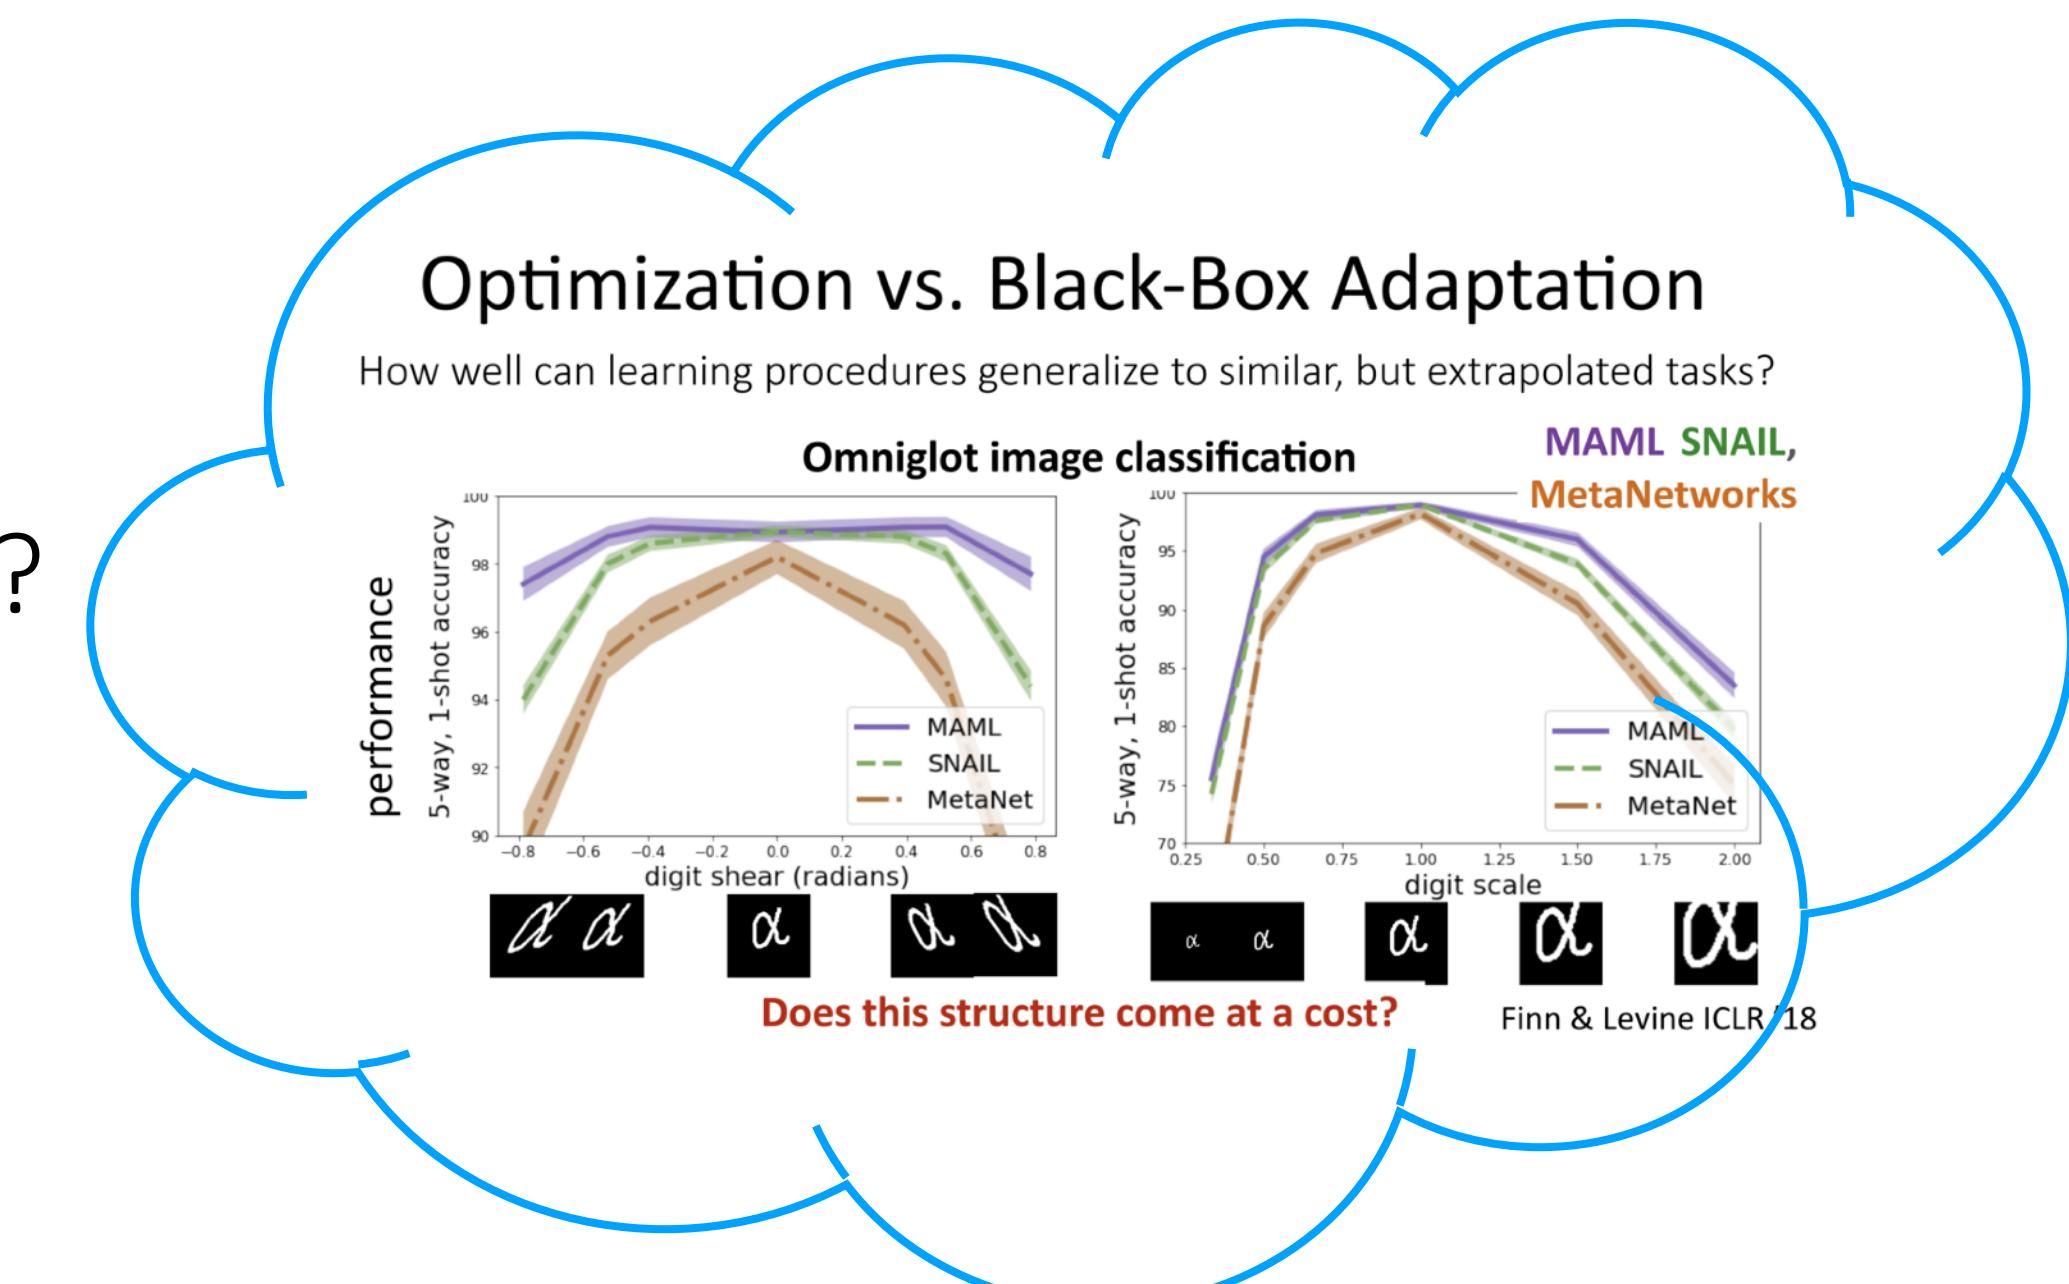
Follow the *meta*-leader (FTML) algorithm:

Store all the data you've seen so far, and **meta-train** on it.

Run update procedure on the current task.

What meta-learning algorithms are well-suited for FTML?

What if $p_t(\mathcal{T})$ is non-stationary?



Online meta-learning experiments

Experiment with **sequences of tasks**:

- Colored, rotated, scaled **MNIST**
- **3D object pose prediction**
- **CIFAR-100** classification

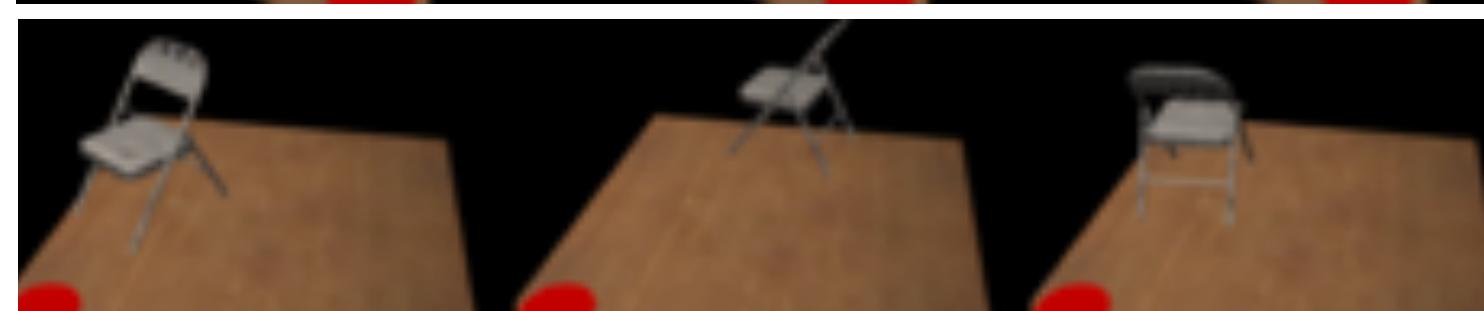
Example pose prediction tasks



plane



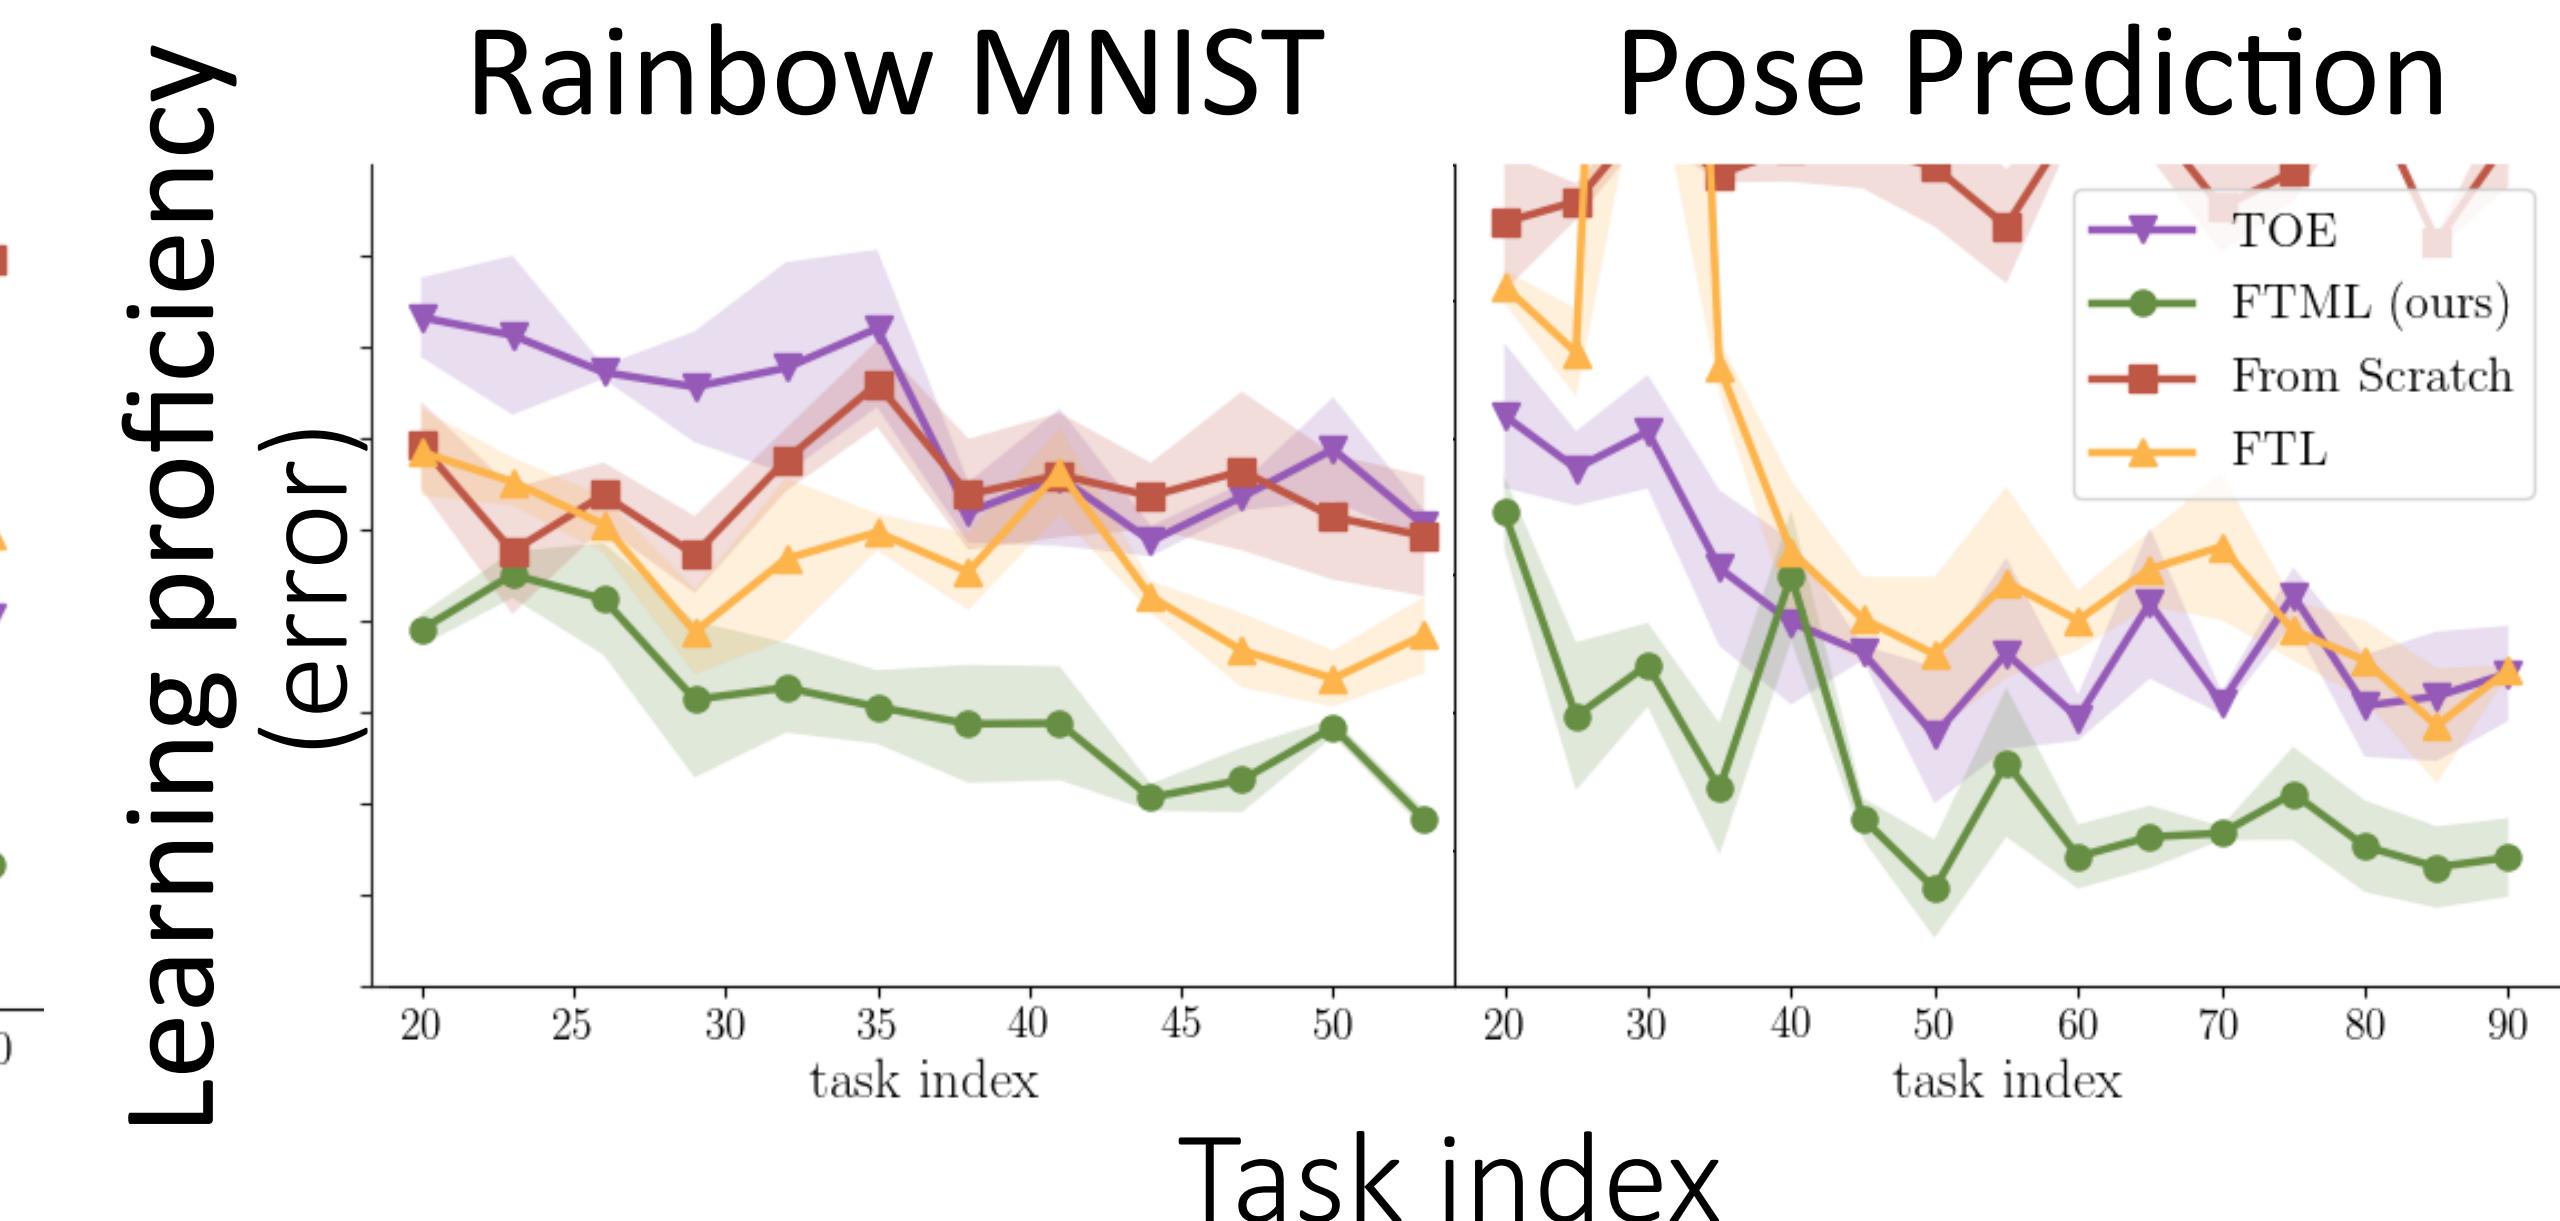
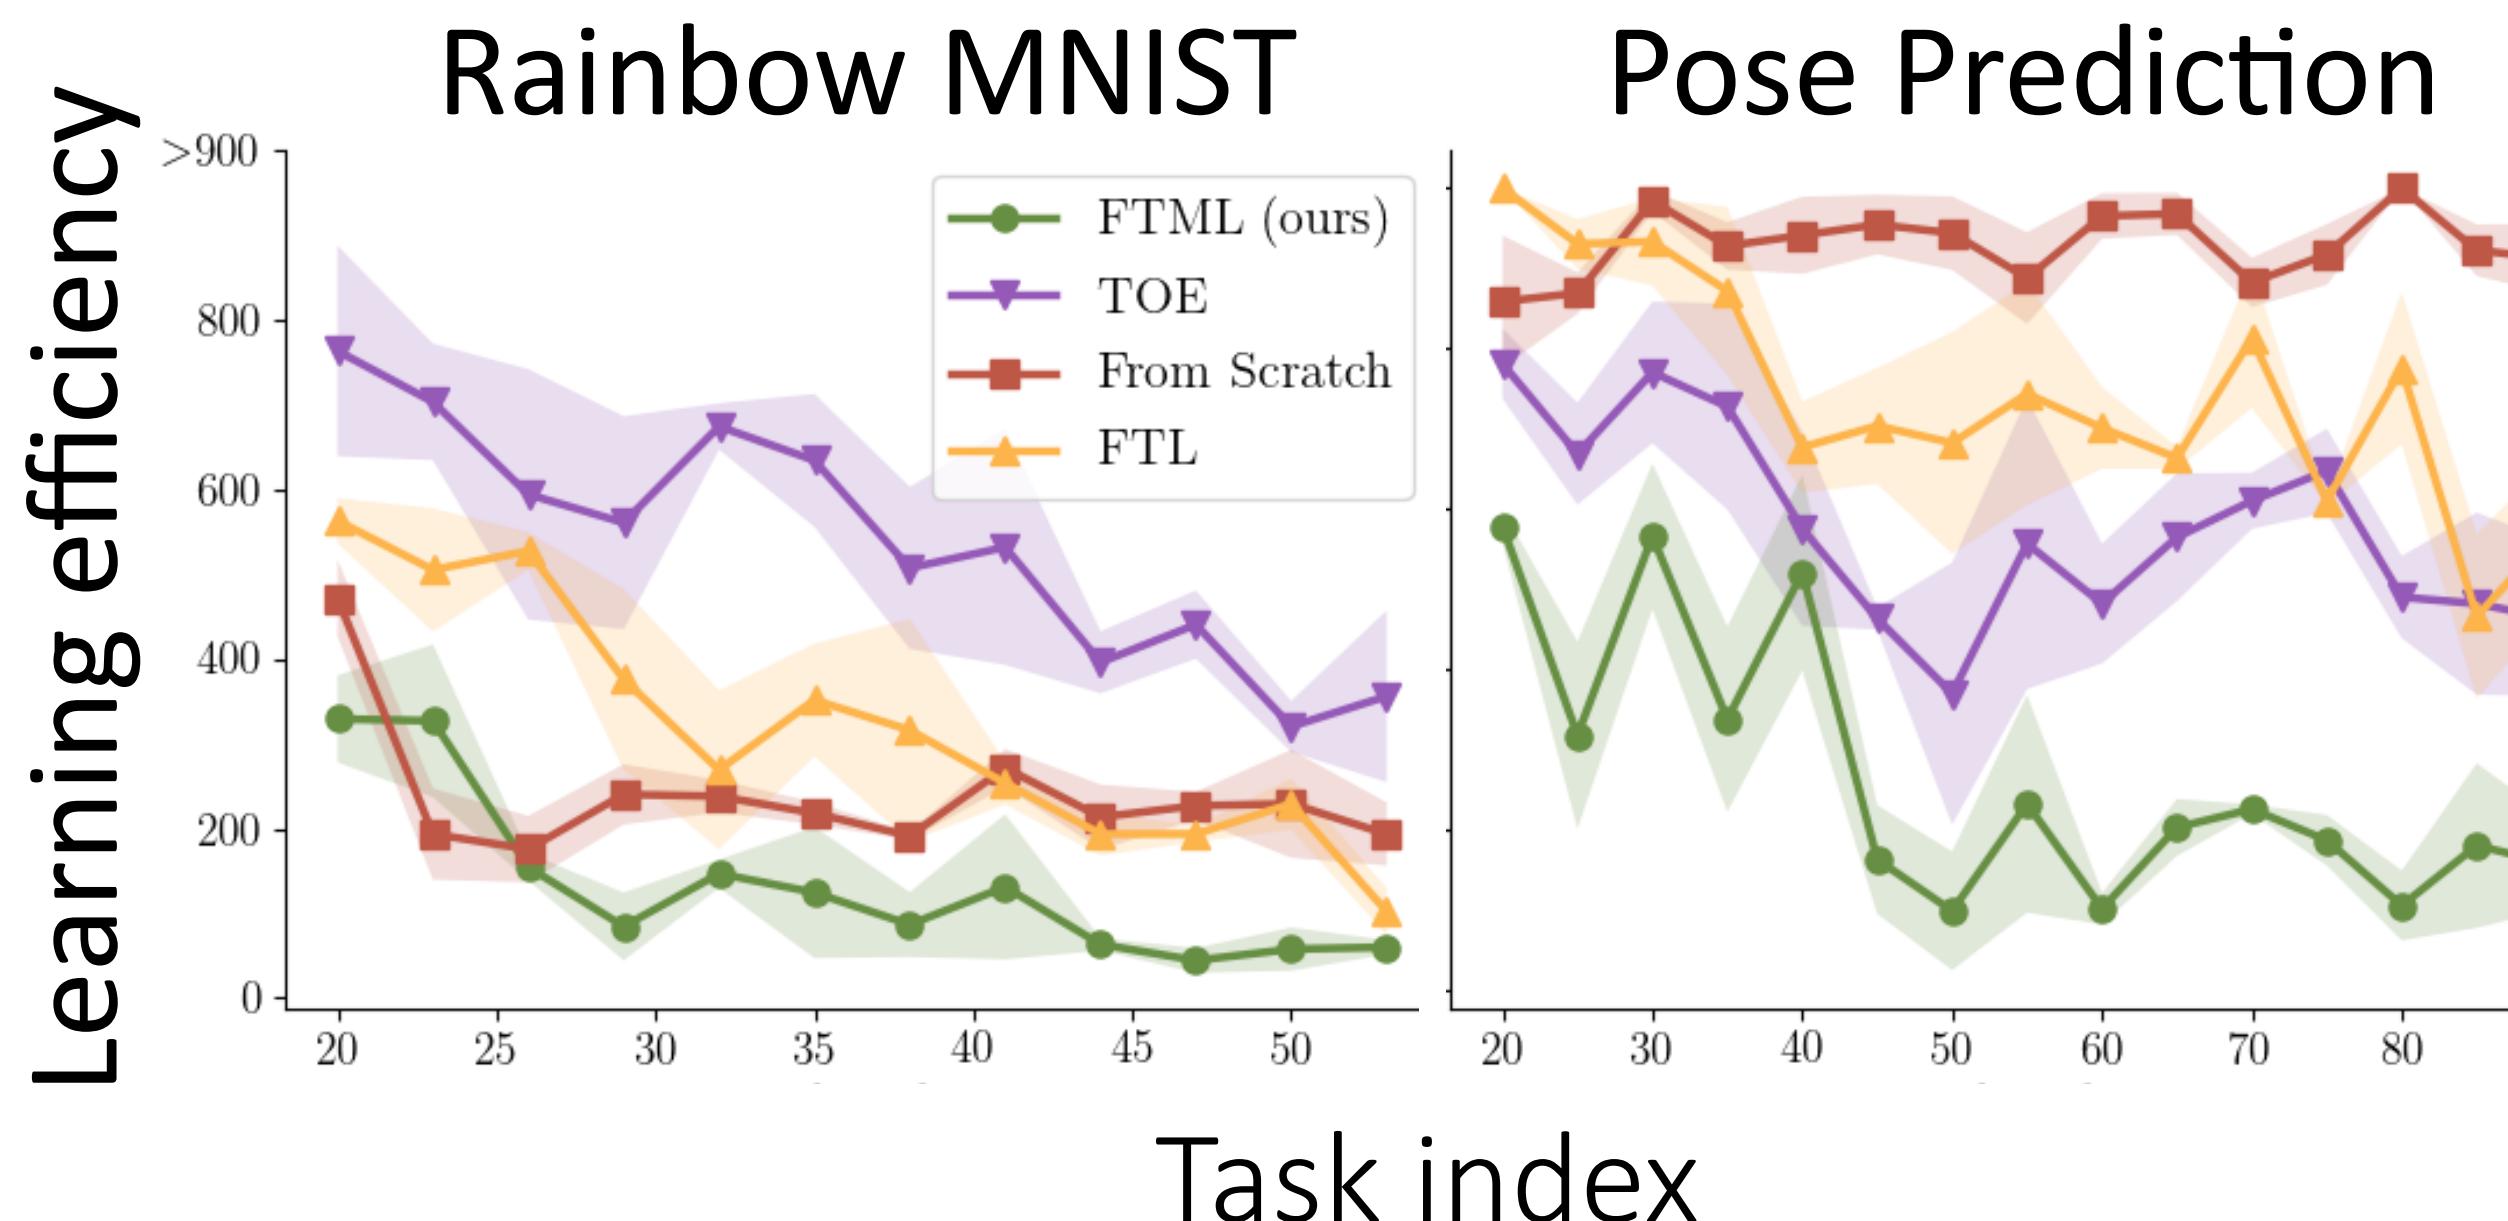
car



chair

Online meta-learning experiments

- Comparisons:
- **TOE (train on everything)**: train on all data so far
 - **FTL (follow the leader)**: train on all data so far, fine-tune on current task
 - **From Scratch**: train from scratch on each task



Follow The Meta-Leader
learns each new task faster & with greater proficiency,
approaches **few-shot learning** regime

Takeaways

Many flavors of lifelong learning, all under the same name.

Defining the problem statement is often the hardest part

Meta-learning can be viewed as a slice of the lifelong learning problem.

A very open area of research.